

Coherence Collapse Analysis

Correlation-Driven Diversity Loss in Complex Coordinating Systems

Eric Moore
eric@ciris.ai

January 2026

Abstract

We apply the Kish design-effect identity $k_{\text{eff}} = k/(1 + \rho(k - 1))$ to coordination among partially-correlated constraints in complex systems, showing that the number of effective independent degrees of freedom collapses to unity as the correlation parameter $\rho \rightarrow 1$. The identity is a Möbius transformation in ρ with asymptotic ceiling $k_{\text{eff}} \rightarrow 1/\rho$ for $k \rightarrow \infty$ at fixed $\rho > 0$: increasing nominal constituent count does not restore diversity lost to correlation. We identify a singularity boundary $K_{\text{req}} \cdot \rho \geq 1$ separating phase regimes (chaos, healthy, rigidity) and derive three closed-form collapse timescales (T_{truth} , T_{entropy} , T_{capture}), together with a stability criterion $\alpha/k \geq d$. These are algebraic results, formalized and machine-verified in Lean 4; the Lean development establishes internal consistency and does not bear on external validity.

We describe a 128-sensor GPU-timing strain gauge instrumenting kernel-jitter correlation on an NVIDIA RTX 4090, and report its characterization: fat-tailed (Student- t) timing statistics, workload-responsive correlation structure, and the domain over which the equicorrelated precondition the identity assumes actually holds. **We do not claim the instrument validates the identity.** $k_{\text{eff}} = k/(1 + \rho(k - 1))$ is an identity; it admits no empirical confirmation, and a comparison of values computed from it against values computed from it returns perfect agreement for any formula whatsoever. Establishing that ρ measured in a live system predicts fragility would require an independently observed measure of effective diversity, which this paper does not supply and which remains open.

We prove an information-theoretic barrier (L-01): the class of marginal-preserving deception patterns undetectable by any method operating on marginal distributions alone is *non-empty*. The theorem establishes existence, not measure; the “~40%” figure appearing in earlier versions follows from a particular illustrative parameterization and is a stated wager, not a result.

CCA is a measurement substrate, not a safety guarantee — it identifies correlation-driven failure modes; domain experts interpret.

Corrections to v3. This version (v5) corrects eighteen defects in v3 (Zenodo 18217688) — the last deposit carrying this paper’s content, v4 having been an accidental upload of a different paper. Several are load-bearing: the stability criterion erred in the permissive direction; the cross-domain and institutional validation sections do not support the conclusions drawn from them; and both results previously offered as hardware validation of the k_{eff} identity are identity checks. Defects, evidence, and dispositions are enumerated in the accompanying corrections note, with a machine witness. Readers comparing versions should start there.

Contents

1	Introduction: The Hypothesis	4
1.1	The Core Claim	4
1.2	The Evidence	4
1.3	The Implications	4

1.4	The Framework	5
1.5	Scope and Falsifiability	5
1.6	Results of conditions 1 and 3	7
1.7	Contributions	9
1.8	Interpretation Guardrails	10
2	Mathematical Foundations	10
2.1	Constraint Geometry	10
2.2	The Defense Function	11
2.3	Stability Condition	12
3	Collapse Timelines	13
3.1	Time to Truth (T_{truth})	13
3.2	Time to Entropy (T_{entropy})	14
3.3	Time to Capture (T_{capture})	14
3.4	Effective Collapse Time	14
4	Phase Space Analysis	15
4.1	The Chaos-Rigidity Spectrum	15
4.2	Classification Algorithm	16
5	Information-Theoretic Limits	17
5.1	The L-01 Barrier	17
5.2	MI Amplification Detection	17
6	Intervention Analysis	18
6.1	Intervention Selection	18
6.2	A Practical Intervention: CIRISAgent	18
7	Validation	19
7.1	Monte Carlo Self-Consistency Check	19
7.2	Formal Verification	20
8	Cross-Domain Validation	21
8.1	Domain Mapping	21
8.2	Empirical Data Sources	21
8.3	k_{eff} Formula Application	21
8.4	Empirical Performance	23
8.4.1	Battery Engine vs NASA Data	23
8.4.2	Institutional Engine vs QoG/Polity	23
8.4.3	Microbiome Engine vs American Gut Project	23
8.5	Testing Rigor: Bug Discovery	23
8.6	Structural Invariant Summary	23
9	Related Work	24
9.1	Safety Engineering: FMEA and Fault Tree Analysis	24
9.2	Control Theory and Stability Analysis	24
9.3	Network Collapse and Phase Transitions	24
9.4	Entropy and Decoherence in Organizations	25
9.5	Byzantine Fault Tolerance and Capture Dynamics	25
9.6	Information-Theoretic Detection Limits	25
9.7	Lineage and Contribution	25

10 Physical Validation: GPU Timing Strain Gauge	26
10.1 Experimental Setup	26
10.2 Kish Identity: Implementation Check (Exp 86)	26
10.3 Software-Induced Collapse (Exp 103) — Withdrawn	28
10.4 Recovery Dynamics (Exp 89)	29
10.5 Leading Indicator Detection (Exp 88)	29
10.6 Signal Architecture	29
10.7 Mean Shift Detection (O-Series Validation, January 2026)	29
10.8 Coherence Collapse Propagation (C-Series, January 2026)	30
10.8.1 C1: k_{eff} Formula Empirical Validation	30
10.8.2 C2: Correlation Propagation Velocity	30
10.8.3 C3: Nucleation Site Distribution	31
10.8.4 C4: Spatial Leading Indicators	31
10.8.5 C-Series Summary	32
10.9 E-Series: Robustness Validation (January 2026)	32
10.10 F-Series: Mechanism and Causality (January 2026)	32
10.11 Implications for CCA	33
11 A Negative Result: The Institutional Case Study	34
11.1 What the engine actually computes	34
11.2 Honest re-scoring	34
11.3 The timing claim is an artifact	35
11.4 What this case study is worth	35
12 Discussion	35
12.1 Limitations	35
12.1.1 Mathematical Simplifications	35
12.1.2 Structural vs Timing Performance	36
12.1.3 Scope Limitations	36
12.2 Outstanding: the validation this paper still lacks	37
12.3 Ethical Considerations	38
13 Conclusion	38
A Notation Reference	40
B Key Formulas	40

1 Introduction: The Hypothesis

1.1 The Core Claim

We propose that complex coordinating systems—whether biological, chemical, institutional, or artificial—share a common failure mode:

Correlation accumulation drives effective diversity toward unity, rendering systems fragile to perturbation regardless of nominal scale.

This claim is formalized through the *effective constraint count*:

$$k_{\text{eff}} = \frac{k}{1 + \rho(k - 1)} \quad (1)$$

where k is the number of constraints (rules, precedents, species, cells) and ρ is their average pairwise correlation. As $\rho \rightarrow 1$, $k_{\text{eff}} \rightarrow 1$ regardless of k . A system with 1,000 highly correlated constraints has the effective diversity of a system with one.

1.2 The Evidence

We applied the variable mapping to three empirical datasets from unrelated domains. Version 3 presented the result as cross-domain validation; that reading is withdrawn, and the corrected status of each row is given below.

Domain	Dataset	Structural Accuracy	Timing Uncertainty
Chemistry	NASA Li-ion (19 cells)	8.1% RMSE	Overestimates final SOH by $\sim 8\%$
Political Science	QoG + Polity V	Withdrawn: 2/5, not 5/5, when scored against a pre-specified outcome	Withdrawn: the “7.6yr early” figure is an initialisation artifact (§11)
Biology	American Gut (2,081 taxa)	Matches AGP norms	FMT dynamics simplified

Table 1: Domain mapping summary, with v3’s institutional row corrected. **The k_{eff} formula is an identity** — it computes effective constraints from measured k and ρ , and admits no empirical confirmation. These rows record how the mapping behaves in each domain; none of them tests the formula.

The same mathematics is *applied to* battery cell degradation, institutional data, and microbiome composition. Version 3 read this as evidence of a structural property common to constraint-based systems. It is not: the shared mathematics is a modelling choice applied in three places, and applying one formula to three datasets demonstrates that the formula can be applied, not that it holds. The institutional application is now reported as a negative result (§11), and the two remaining rows carry proxy ρ values rather than measured intraclass correlations (Table 8).

1.3 The Implications

If this failure mode generalizes beyond the three tested domains, it may have broader consequences:

1. **The failure is invisible:** Systems accumulate correlation while appearing healthy (k grows). The collapse of k_{eff} is not directly observable without measuring ρ .
2. **The failure tends toward simpler configurations:** Under information-theoretic interpretation, correlated systems require less information to describe than diverse ones. Without active maintenance of diversity, systems drift toward homogeneity. (Note: This is an information-theoretic statement, not a claim about physical thermodynamics.)
3. **AI accelerates the failure asymmetrically:** AI systems increase constraint count ($k \uparrow$) and correlation ($\rho \uparrow\uparrow$) simultaneously, while externalizing sustainability costs ($\sigma \downarrow$). This masks diversity collapse behind apparent scale [39, 41, 42].

As a metaphor (not a probabilistic claim), correlation accumulation shares structural features with proposed “Great Filter” mechanisms: it operates invisibly, scales with technological capability, and resists intervention once advanced. This analogy is illustrative, not predictive [40, 45].

1.4 The Framework

Coherence Collapse Analysis (CCA) is the formal apparatus for analyzing this failure mode. It synthesizes techniques from:

- Failure Mode and Effects Analysis (FMEA) in safety engineering
- Stability analysis in control theory
- Phase transition analysis in statistical mechanics
- Information-theoretic bounds from detection theory

CCA provides:

1. **Failure modes:** How coherence is lost (deception, entropy, capture)
2. **Attractor states:** Where systems tend to drift (chaos vs. rigidity)
3. **Intervention windows:** When corrective action remains effective
4. **Detection limits:** What can and cannot be observed

1.5 Scope and Falsifiability

CCA Is	CCA Is Not
Structural analysis of failure conditions	Prediction of historical outcomes
Identification of phase boundaries	Fortune-telling or prophecy
An identity applied across domains	Cross-domain validated mathematics
Engineering-style failure analysis	Social physics or psychohistory
A measurement substrate	A safety guarantee

Table 2: Scope boundaries for Coherence Collapse Analysis

Falsification conditions. Version 3 stated three conditions, none of which could fire: that “the k_{eff} formula fails in additional domains” (it is an identity, and cannot fail), that “systems with high ρ demonstrate sustained resilience” (absorbed by this paper’s own block-structure caveat, which converts any such observation into confirmation), and that “correlation accumulation reverses spontaneously” (no timescale, magnitude, or measurement procedure given). They are replaced with conditions carrying an observable, a threshold, and a procedure.

1. **Effective-count disagreement.** Compute $N_{\text{eff}} = k/\lambda_{\text{max}}$ from the largest eigenvalue of a measured correlation matrix and $k_{\text{eff}} = k/(1 + \bar{\rho}(k - 1))$ from the same data. These

are distinct functionals over distinct inputs. If they diverge by more than 25% across the operating range of ρ on a substrate where the equicorrelated precondition holds (condition 2), the identity is not the right effective-count law for that substrate. *Run: does not fire*, at 20.4% median disagreement (§1.6).

2. **Precondition failure in the operating band.** The Kish form assumes exchangeability. If the spectral entropy deficit departs from its closed form by more than 10% anywhere in the band where a deployment actually operates, k_{eff} figures from that band are not licensed. *This condition has already fired on the GPU substrate below $\rho \approx 0.04$, and bounds the domain of every k_{eff} value reported here.*
3. **No two-sided corridor.** Measure fragility across a ρ range spanning both corridor edges on any substrate. If fragility is monotone in ρ rather than U-shaped — in particular if it does not rise again above the upper edge — the corridor claim fails for that substrate. *Run twice under two fragility measures, which disagree in direction: **UNRESOLVED** (§1.6).*
4. **Volume decay under correlation.** The decay law $V(k) = V_0 e^{-\lambda k_{\text{eff}}}$ predicts a specific effective count from measured volume reduction. If simulated or measured decay implies an effective count differing from k_{eff} by more than a factor of 1.5, the law is wrong. *This condition has already fired: the implied count exceeds k_{eff} by $1.9\times$ to $4.1\times$ at $\rho = 0.5$, growing with k (§12.2).*
5. **Threshold transfer.** The collapse rule $\sigma < \sigma_{\min}$ or $f > f_{\max}$ carries two free parameters, fitted per domain. Fix one pair on one domain and apply it unchanged to another. If classification accuracy does not exceed a base-rate-matched null under a pre-specified outcome definition, the rule is a per-domain fitted classifier rather than a structural invariant. *Currently unmet: no threshold pair in this paper has been shown to transfer, and the one domain scored against a pre-specified outcome (§11) performs below chance.*
6. **Sensitivity calibration.** The detection model $P = 1 - e^{-\beta \varepsilon}$ has a free β , never independently calibrated. Estimate β on one substrate and predict detection rates on a second. If β must be refit per substrate, the model has no predictive content beyond curve-fitting, and quantities derived from a chosen β — including the “ $\sim 40\%$ ” undetectable fraction of earlier versions — carry none either. *Currently unmet: $\beta = 10$ is a stipulation, and §7.1 shows the only test of it samples from the model being tested.*

Conditions owed by the fires. A fired condition that leaves no successor lets the repaired claim return unchallenged. Three more are therefore entered, each generated by a specific result above.

7. **Replacement decay law.** Condition 4 fired, so $V = V_0 e^{-\lambda k_{\text{eff}}}$ is not the correct law. Any proposed replacement must (a) reduce to k at $\rho = 0$, (b) reproduce the measured over-prediction of effective count, $1.9\times$ to $4.1\times$ at $\rho = 0.5$ growing in k , and (c) account for the degradation observed at large k even at $\rho = 0$ (implied count 25.7 against $k = 32$). A law failing any of the three is not the repair, whatever else it fits.
8. **Sign of the effective-count residual.** Condition 1 passed at 20.4%, but $k/\lambda_{\max} \leq k_{\text{eff}}$ in 21 of 21 matrices — a one-directional bias, not noise. A corrected estimator must eliminate the residual’s *sign*, not merely shrink its magnitude. A proposal that halves the disagreement while leaving it systematically signed has fitted the size of the error without finding its mechanism.
9. **Prior validation of fragility measures.** Any measure used to support a corridor claim must be shown k_{eff} -independent *before* it is used, on data spanning the corridor. Three measures have now failed this check after the fact (§1.6). Post-hoc justification of a fragility measure is not admissible for corridor claims in this framework.

Of nine conditions: two fired, three are unmet, one is unresolved, one passes, and three are newly entered by the fires. That distribution is reported rather than smoothed, because a condition that cannot fire is not a falsification condition, and the three carried in version 3 could not. They are withdrawn rather than reworded.

1.6 Results of conditions 1 and 3

Both were runnable on data and hardware already in hand, so they were run rather than left standing. Pre-registered before either quantity was computed (hash and predictions in replication/).

Condition 1 — does not fire, but the residual is systematic. Over the 21 stored C1 correlation matrices ($k = 16$), median disagreement between $k/\lambda_{\max}$ and k_{eff} is **20.4%** on the precondition-satisfying set ($\rho \geq 0.04$), against a 25% bar. **Kish passes this test** — the first non-circular check of the identity anywhere in this corpus, and the measurement §10.2 says the paper lacks.

Two features of the residual matter more than the pass. First, the direction is not random: $k/\lambda_{\max} \leq k_{\text{eff}}$ in **21 of 21** matrices, so the identity systematically *over-credits* effective diversity relative to the spectrum — the unsafe direction, since a system reading k_{eff} believes it has more independent constraints than the eigenstructure supports. Second, the disagreement concentrates where the precondition is weakest: median 23.2% for $\rho < 0.10$ against 11.1% for $\rho \geq 0.20$, consistent with the Exp 117 finding that exchangeability fails at low correlation. A 20.4% median against a 25% bar is a pass, not a comfortable one.

Condition 3 — fires. The rigidity arm is not observed. The rigidity regime was reached for the first time on this substrate (ρ up to 0.868, against the F-series maximum of 0.171). Perturbation response, the F-series fragility measure:

Regime	n	mean response	sd
healthy ($0.10 \leq \rho < 0.43$)	7	36.7%	17.5%
rigidity ($\rho \geq 0.43$)	5	10.0%	5.2%

Fragility does not rise above the upper edge — it *falls*, by a factor of 3.7. Under the pre-registered rule the condition fires: **the corridor is not two-sided on this substrate**. Combined with Remark 10.5, the position is that the chaos arm is supported and the rigidity arm is now measured and absent.

That firing does not stand, and the reason is instructive. Perturbation response conflates two distinct things: a genuinely robust system, and a sensor array that has stopped discriminating. At $\rho \rightarrow 1$ the array approaches a single effective sensor, and a single sensor responds less to spatial perturbation whether or not it is robust. So the falling response is equally consistent with “rigidity is not fragile here” and with “the fragility measure stops working exactly where the corridor needs it.” The F-series used this same measure for the chaos arm, so the objection reaches that result too.

Condition 3, re-run under recovery time — unresolved. The objection was addressed by building the measure it calls for. Recovery time τ is estimated from the AR(1) lag-1 autocorrelation of the array mean after a sustained perturbation is released — the standard critical-slowness statistic, and a *temporal* property, so it does not degrade as the array loses spatial discrimination. Sampling is synchronous in every condition (time-resolved array state requires it), and ρ is varied by memory sharing rather than by stream asynchrony, so the correlation control is independent of the sampling method.

Regime	n	τ (array)	sd	$\tau_{\text{array}}/\tau_{\text{single}}$
healthy ($0.10 \leq \rho < 0.43$)	4	2.21	1.17	2.15
rigidity ($\rho \geq 0.43$)	12	3.67	2.82	1.32

The two measures disagree in direction. Recovery time is *longer* at rigidity (+1.46), the sign critical slowing down predicts and the opposite of what response amplitude showed — which is what one expects if the first result was carrying the confound. But the difference is smaller than the pooled standard deviation (2.29), so it is not significant; the healthy bin holds only 4 of 16 runs because achieved ρ landed high; and the pre-registered control failed: the array-to-single-sensor τ ratio reached 2.15 in the healthy bin against an admissibility bound of 2.0, so this measure cannot be certified free of its own k_{eff} dependence either.

Under the pre-registered decision rule a failed control makes the condition **UNRESOLVED**, and neither the earlier firing nor a survival may be claimed. **The version-3 corridor claim therefore remains unsupported on this substrate, and the first run’s apparent refutation of it is withdrawn as well.** What is established is narrower and worth stating on its own: the rigidity regime is reachable here (ρ to 0.868), two defensible fragility measures give opposite signs across the upper edge, and neither is clean.

Condition 3, third attempt — and why it cannot be settled on this instrument. Detection latency at fixed false-positive rate was the remaining candidate, and it was built: the array is a detector, so fragility is degraded detection, with the decision threshold calibrated per condition on that condition’s *own* baseline to a 5% false-positive rate. That calibration was the point — it normalises away the variance scale that contaminated both earlier measures.

It failed, for a reason worth recording. At *zero* perturbation the measure reports a latency of 3.3 samples where its own 5% calibration implies roughly 20: the array’s baseline drifts between calibration and test, so the detector is timing drift rather than signal. Perturbation response is also a cliff rather than a gradient — latency sits near the false-alarm floor across a range of drive strengths and then collapses to 1.0 — leaving no dynamic range in which conditions could be compared.

The three failures share one cause. **Every k_{eff} -independent fragility measure requires the array state at each time index, which requires synchronous sampling across all sensors; and synchronous sampling on this instrument is itself a correlating operation, holding ρ at or above 0.45 in every condition attempted.** Varying workload heterogeneity to break the coupling raised ρ further ($0.468 \rightarrow 0.713$) rather than lowering it. The healthy band is therefore not reachable by any time-resolved measure on this hardware, and the comparison the corridor claim needs cannot be constructed here.

F-3 is settled as unsettlable on this substrate, which is a property of the instrument rather than of the corridor. Deciding it requires an instrument in which correlation varies independently of sampling synchrony — physically separated sensors across multiple devices, or asynchronous per-sensor sampling interpolated onto a common time grid. Until then the corridor’s rigidity arm is neither supported nor refuted, and this paper claims neither.

The chaos arm was not re-tested under either measure: no run reached $\rho < 0.10$, so condition 3’s chaos-side prediction is untested rather than replicated.

Two claims that appeared in version 3’s “structural invariants” receive no falsification handle here, because none can exist for them. $k_{\text{eff}} = k/(1 + \rho(k - 1))$ is an **identity**; $f = 1 - \sigma$ is a **definition** in the engines that report it as validated. Neither can be contradicted by any measurement, and the appropriate response is to stop describing them as invariants that were checked — which this version does — rather than to invent a falsifier they cannot have.

1.7 Contributions

Analytic results.

1. The **Möbius structure** of the effective-count map: k_{eff} is strictly monotone decreasing in ρ and saturates at $1/\rho$ as $k \rightarrow \infty$ (§2.2). Scale cannot restore diversity that correlation has already collapsed — at $\rho = 0.1$ the ceiling is 10 effective constraints, at $\rho = 0.5$ it is 2, however many are added. This is substrate-independent and requires no calibration.
2. A **stability condition** $\alpha/k \geq d$ separating growing from decaying systems, with the corollary that constant-generation systems are eventually unstable for any decay rate.
3. Derivation of **three collapse timescales** (T_{truth} , T_{entropy} , T_{capture}) in closed form, and a singularity boundary $K_{\text{req}} \cdot \rho \geq 1$ beyond which added constraints cannot reduce incoherence.
4. A **chaos–rigidity phase decomposition** and the corridor between them. The chaos arm is measured on the GPU substrate; the rigidity arm is not (Remark 10.5).
5. An information-theoretic **existence result** (L-01): the class of marginal-preserving incoherence patterns undetectable from marginals alone is non-empty. Existence only — the theorem gives no measure over that class.
6. Formal verification of the algebraic results in **Lean 4**, establishing internal consistency and, as §7.2 states, nothing about external validity.

Instrument and measurements.

7. A **128-sensor GPU-timing strain gauge**, characterized: fat-tailed (Student- t) timing statistics, workload-responsive correlation structure, first measurement of correlation propagation velocity (0.5 m/s), and a bounded domain over which the equicorrelated precondition the identity assumes actually holds (Exp 117: within 0.8% at $\rho = 0.394$, diverging to 70.8% at $\rho \approx 0.03$).
8. A **measurement protocol** for correlation on this class of instrument, derived from a pre-registered replication: a shuffle null (at $n = 30$ samples the null’s 95th percentile is $\rho \approx 0.15$), repeated trials with intervals rather than single draws, and rejection of the first measurement after idle. Applying it withdrew one of this paper’s own headline results.
9. An **open-source implementation** — Python modules, Lean proofs, simulation scripts, pre-registrations, and the machine witness for every algebraic correction in this version (github.com/CIRISAI/RATCHET).

What this version does not contribute. Version 3 of this paper claimed cross-domain empirical validation of the k_{eff} formula and hardware confirmation of its central mathematical claim. **Both are withdrawn.** k_{eff} is an identity: it admits no empirical confirmation, and the two results previously offered as confirming it computed the identity and compared it against itself. The institutional application does not test k_{eff} at all and is reported here as a negative result (§11). Eighteen defects are enumerated with evidence and dispositions in the accompanying corrections note; the two that resist repair — the k_{eff} substitution in Theorem 2.1, and the deceptive/honest asymmetry — are stated as open in §12.2 with measurements of how far each fails.

1.8 Interpretation Guardrails

To prevent misreading, we state explicitly what CCA does **not** claim:

- **No inevitability:** Correlation accumulation is a tendency, not a fate. Systems can and do maintain diversity through active intervention.
- **No timeline predictions:** CCA identifies *structural fragility*, not when collapse will occur. Version 3 cited a “7.6-year timing error” here as illustrating the limitation; that figure is itself withdrawn as an initialisation artifact (§11), and the limitation stands without it — no lead-time claim in this paper survives comparison against a null that fires at t_0 .
- **No macro-outcome claims:** We do not predict civilizational collapse, AI doom scenarios, or specific historical events. CCA is engineering risk analysis, not prophecy.
- **No causal claim:** CCA does not claim correlation *causes* collapse—only that high correlation is a sufficient structural bottleneck for fragility, regardless of underlying mechanism (see F4: common cause confirmation in Section 10).
- **“Great Filter” is metaphorical:** The term is used as an illustrative analogy for correlation-driven failure, not as a probabilistic claim about the Fermi paradox.

The framework’s value lies in making hidden fragility *legible*—identifying when systems have drifted into dangerous phase space, not in predicting what happens next.

2 Mathematical Foundations

2.1 Constraint Geometry

The foundation of CCA is the observation that system coherence is maintained through *constraints*—rules, precedents, norms, or structural features that limit the space of possible behaviors.

Definition 2.1 (Constraint Space). A **constraint space** is a tuple $(V, \mathcal{C}, \rho)$ where:

- $V \subseteq \mathbb{R}^n$ is the space of possible system states
- $\mathcal{C} = \{C_1, \dots, C_k\}$ is a set of constraint halfspaces
- $\rho : \mathcal{C} \times \mathcal{C} \rightarrow [0, 1]$ is the pairwise correlation function

The *feasible region* is the intersection $\bigcap_{i=1}^k C_i$, representing states consistent with all constraints. Deceptive or incoherent behaviors correspond to states outside this region.

Definition 2.2 (Effective Constraint Count). Given k constraints with average pairwise correlation ρ , the **effective constraint count** is:

$$k_{\text{eff}} = \frac{k}{1 + \rho(k - 1)} \quad (2)$$

Remark 2.1 (Statistical Provenance). This formula is mathematically identical to the **Kish design effect** for effective sample size in survey statistics, which accounts for autocorrelation in clustered samples. We do not claim novelty for the formula itself—only for its application to constraint-based system analysis. The “validation” of k_{eff} across domains confirms that the algebra of effective sample sizes holds when variables are appropriately mapped; it does not establish that such mappings are causally meaningful in all contexts.

This captures the intuition that correlated constraints provide redundant information:

- When $\rho = 0$ (independent): $k_{\text{eff}} = k$
- When $\rho \rightarrow 1$ (fully correlated): $k_{\text{eff}} \rightarrow 1$

Remark 2.2 (Möbius Structure and Asymptotic Ceiling). For fixed k , the map $\rho \mapsto k_{\text{eff}}(k, \rho)$ is the Möbius transformation with coefficients $(a, b, c, d) = (0, k, k - 1, 1)$. The partial derivative $\partial k_{\text{eff}} / \partial \rho = -k(k - 1) / (1 + \rho(k - 1))^2 < 0$ for $k > 1$, so k_{eff} is strictly monotone decreasing in ρ on $[0, 1]$. Taking $k \rightarrow \infty$ at fixed $\rho > 0$:

$$\lim_{k \rightarrow \infty} k_{\text{eff}}(k, \rho) = \lim_{k \rightarrow \infty} \frac{k}{1 + \rho(k - 1)} = \frac{1}{\rho}.$$

Effective dimensionality saturates at the inverse correlation regardless of nominal constituent count. **Adding more correlated constituents cannot restore diversity that correlation has already collapsed:** at $\rho = 0.1$ the ceiling is $k_{\text{eff}} = 10$; at $\rho = 0.5$ it is $k_{\text{eff}} = 2$; at $\rho = 0.9$ it is $k_{\text{eff}} \approx 1.11$. This is the algebraic core of the “scale alone cannot compensate for lost diversity” claim, and it is independent of any substrate-specific calibration.

Remark 2.3 (Scalar ρ Limitation). The average pairwise correlation ρ is a scalar summary that can obscure important structure. **High average ρ does not imply collapse when correlation is modular rather than global.** A system with two independent clusters, each internally correlated, may show high average ρ but retain $k_{\text{eff}} = 2$ (the number of independent modules). The E-series experiments (Section 10) explicitly test for block-diagonal correlation structure. Future work should extend CCA to use correlation tensors that preserve this information.

Theorem 2.1 (Volume Decay). Under the following assumptions:

- (a) Constraints are drawn i.i.d. from a distribution over halfspaces with bounded support
- (b) The initial feasible region V_0 is bounded and convex
- (c) Constraint normals have mean zero and covariance Σ with $\|\Sigma\| \leq M$ for some $M > 0$

the volume of the feasible region decays exponentially with effective constraints:

$$V(k) = V_0 \cdot e^{-\lambda \cdot k_{\text{eff}}} \cdot (1 + O(1/\sqrt{k_{\text{eff}}})) \quad (3)$$

where $\lambda > 0$ is the decay constant determined by constraint geometry, and the error term vanishes as $k_{\text{eff}} \rightarrow \infty$.

Proof sketch. Under assumption (a), each independent constraint removes a fraction of the remaining volume. By Grünbaum’s theorem on log-concavity of volume under halfspace intersection, the reduction is multiplicative. For k_{eff} effective (non-redundant) constraints, this yields $V(k_{\text{eff}}) \approx V_0 \cdot \prod_{i=1}^{k_{\text{eff}}} (1 - \delta_i)$ where δ_i is the fraction removed by constraint i . Taking logs and applying the law of large numbers under assumption (c), $\sum \ln(1 - \delta_i) \approx -\lambda \cdot k_{\text{eff}}$ for some λ depending on the constraint distribution. The correlation adjustment (k_{eff} vs k) follows from the effective sample size formula: correlated constraints provide redundant information, reducing the effective count by factor $(1 + \rho(k - 1))^{-1}$. $\square$

Remark 2.4 (Applicability). The exponential form is an approximation valid when constraints are “generic” (no degenerate intersections) and sufficiently numerous. For small k or highly structured constraint sets, empirical validation is required.

2.2 The Defense Function

Definition 2.3 (Defense Function). The **defense function** J quantifies the computational cost an adversary must pay to maintain incoherent behavior:

$$J = k_{\text{eff}} \cdot \lambda \cdot \sigma \quad (4)$$

where:

- k_{eff} : Effective scale — constraint count *after* the correlation discount

- λ : Strictness (constraint strength)
- σ : Sustainability (maintenance capacity)

Remark 2.5 (Correction: diversity is carried by k_{eff} , not a separate factor). Version 3 of this paper defined $J = k_{\text{eff}} \cdot (1 - \rho) \cdot \lambda \cdot \sigma$. That form **double-counted correlation**: $k_{\text{eff}} = k/(1 + \rho(k - 1))$ already contains the diversity discount, so multiplying by a further $(1 - \rho)$ applies it twice and drives $J \rightarrow 0$ as $\rho \rightarrow 1$. The latter contradicts the single-constraint floor — a fully correlated system is no safer than *one* constraint, but it is not *worse* than one. The corrected three-factor form above is the one carried by the CIRIS Constitution (CC 6.2.2). Diversity enters J **only** through k_{eff} . See `CORRECTIONS_v3.md`.

Remark 2.6 (Definitional Status of J). The defense function J is a **modeling choice**, not a derived quantity. We do not claim that J is the unique or correct measure of system resilience—only that this particular combination of terms produces useful analysis across multiple domains. Alternative formulations (e.g., weighted products, nonlinear combinations) may be appropriate for specific applications.

This formulation exhibits a conceptual duality: systems with high defense J against incoherence also tend to have high capacity for coordinated flourishing. High-defense systems are also high-capacity systems.

Proposition 2.2 (J-C Duality Interpretation). Let C denote an abstract “capacity for coordinated flourishing.” If we operationalize C as the product of:

- effective coordination channels (k_{eff} , which already discounts channels that are not independent),
- enforcement strength (λ), and
- resource sustainability (σ),

then $J = C$ by definition. This is not a derived result but a **definitional choice** that interprets defense and capacity as dual aspects of the same underlying quantity.

Remark 2.7. The J-C duality is interpretive rather than mathematical. It asserts that the same structural features that make a system hard to subvert also make it capable of positive coordination. This claim is falsifiable: if empirical systems show high J with low flourishing capacity (or vice versa), the interpretation fails.

The defense function provides a single scalar that summarizes system health. Its dynamics determine whether the system is growing stronger or collapsing.

2.3 Stability Condition

Definition 2.4 (Collapse Rate). The **collapse rate** is the time derivative of the defense function:

$$R_c = \frac{dJ}{dt} \quad (5)$$

Theorem 2.3 (Stability Condition). A system is stable if and only if $R_c \geq 0$. Under the simplifying assumption that λ is constant, this reduces to:

$$\text{Stable} \iff \frac{\alpha}{k} \geq d \quad (6)$$

where $\alpha = dk/dt$ is the constraint generation rate and d is the sustainability decay rate. Note that the criterion is governed by the **nominal** count k , not by k_{eff} .

Derivation. From $J = k_{\text{eff}} \cdot \lambda \cdot \sigma$, apply the product rule:

$$\frac{dJ}{dt} = \lambda \left[\frac{dk_{\text{eff}}}{dt} \sigma + k_{\text{eff}} \frac{d\sigma}{dt} \right] \quad (7)$$

Assume the dominant dynamics are:

- (i) k_{eff} grows proportionally to new constraints: $dk_{\text{eff}}/dt \approx \alpha/(1 + \rho(k - 1))$
- (ii) ρ changes slowly: $d\rho/dt \approx 0$ (quasi-static)
- (iii) σ decays exponentially without input: $d\sigma/dt = -d \cdot \sigma$

Substituting these together with $k_{\text{eff}} = k/(1 + \rho(k - 1))$:

$$R_c \approx \lambda\sigma \left[\frac{\alpha}{1 + \rho(k - 1)} - d \cdot \frac{k}{1 + \rho(k - 1)} \right] = \frac{\lambda\sigma(\alpha - dk)}{1 + \rho(k - 1)} \quad (8)$$

Since $\lambda, \sigma > 0$ and $1 + \rho(k - 1) > 0$ for $\rho \in [0, 1]$, $k \geq 1$, the sign of R_c is governed entirely by the factor $(\alpha - dk)$:

$$R_c \geq 0 \iff \alpha \geq d \cdot k \iff \frac{\alpha}{k} \geq d \quad (9)$$

The inequality is strict for stability margin. $\square$

Remark 2.8 (Correction to v3: the criterion is α/k , not α/k_{eff}). Version 3 of this paper stated the criterion as $\alpha/k_{\text{eff}} > d$. That does not follow from the derivation above: the common denominator $1 + \rho(k - 1)$ cancels between the two terms, leaving $(\alpha - dk)$ and hence $\alpha/k \geq d$. Because $k_{\text{eff}} \leq k$, the v3 criterion is **too permissive by exactly** $1 + \rho(k - 1)$ — the design-effect factor this paper exists to introduce — and it errs in the unsafe direction, certifying as stable systems whose defense function is strictly decreasing. For example, at $k = 100$, $\rho = 0.5$, $d = 0.1$, $\alpha = 5.099$: the v3 criterion gives $\alpha/k_{\text{eff}} = 2.575 > d$ (“stable”) while $dJ/dt < 0$. The correction is independent of the J -form correction above: with or without the $(1 - \rho)$ factor, the boundary is $\alpha = dk$. Machine witness: `corrections_witness.py`; see `CORRECTIONS_v3.md`.

Corollary 2.4 (Static Systems Are Doomed). If α is constant and k grows without bound, the ratio α/k decreases monotonically toward zero. Eventually $\alpha/k < d$, violating stability.

Remark 2.9 (Why the corollary requires k rather than k_{eff}). This corollary is **false** under the v3 criterion and is restored only by the correction above. The v3 argument asserted that k growing implies k_{eff} growing, so that α/k_{eff} falls below d . But k_{eff} does *not* grow without bound: it saturates at $1/\rho$, so α/k_{eff} converges to $\alpha\rho$ *from above* rather than decaying. Whenever $\alpha\rho > d$ the v3 criterion is never violated, no matter how large k becomes — at $\rho = 0.5$, $d = 0.1$, $\alpha = 1$ the ratio floors at 0.5, permanently above d . Under the corrected criterion the conclusion holds unconditionally, since $\alpha/k \rightarrow 0$ for any fixed α . The headline claim survives; it does not survive the reasoning as published.

Remark 2.10 (Dimensional Analysis). The stability condition is dimensionally consistent: α has units [constraints/time], k is dimensionless, and d has units [1/time]. The ratio α/k gives the rate of constraint generation per existing constraint, which must exceed the decay rate.

This is a key insight: static systems cannot maintain coherence indefinitely. The stability condition demands continual renewal.

3 Collapse Timelines

CCA identifies three primary collapse modes, each with a characteristic timeline.

3.1 Time to Truth (T_{truth})

Definition 3.1 (Required Constraints). Given initial volume V_0 , target safety threshold ε , and decay constant λ , the required effective constraints are:

$$K_{\text{req}} = \frac{-\ln(\varepsilon/V_0)}{\lambda} \quad (10)$$

Definition 3.2 (Critical Correlation). The **critical correlation threshold** is:

$$\rho_{\text{crit}} = \frac{1}{K_{\text{req}}} \quad (11)$$

If $\rho \geq \rho_{\text{crit}}$, the system cannot generate sufficient effective constraints regardless of scale.

Theorem 3.1 (Time to Truth). The time until incoherent behavior becomes computationally prohibitive is:

$$T_{\text{truth}} = \frac{K_{\text{req}}(1 - \rho)}{\alpha(1 - K_{\text{req}} \cdot \rho)} \quad (12)$$

Theorem 3.2 (Singularity Condition). When $K_{\text{req}} \cdot \rho \geq 1$:

$$T_{\text{truth}} \rightarrow \infty \quad (13)$$

This is the **rigidity boundary**—an “echo chamber” where correlated constraints provide no additional security regardless of scale or time.

3.2 Time to Entropy (T_{entropy})

Systems require ongoing maintenance to prevent entropic decay.

Definition 3.3 (Sustainability Dynamics). The sustainability integral evolves as:

$$\sigma(t) = \sigma_0 \cdot e^{-d \cdot t} + \int_0^t w \cdot S(\tau) \cdot e^{-d(t-\tau)} d\tau \quad (14)$$

where $S(t)$ is the signal (value production) and w is the weight.

Theorem 3.3 (Time to Entropy). The time until sustainability falls below the revocation threshold σ_{min} is:

$$T_{\text{entropy}} = \frac{\ln(\sigma/\sigma_{\text{min}})}{d} \quad (15)$$

This represents the “black hole” failure mode: systems that consume resources without producing value eventually collapse.

3.3 Time to Capture (T_{capture})

Distributed systems can be captured through gradual corruption of nodes.

Theorem 3.4 (Time to Capture). For a federation with n nodes, f currently compromised, and capture rate r_{cap} :

$$T_{\text{capture}} = \frac{(n/3) - f}{r_{\text{cap}}} \quad (16)$$

where $n/3$ is the Byzantine fault tolerance threshold.

3.4 Effective Collapse Time

Definition 3.4 (Effective Collapse Time). The **effective collapse time** is the minimum of all collapse modes:

$$T_{\text{eff}} = \min(T_{\text{truth}}, T_{\text{entropy}}, T_{\text{capture}}) \quad (17)$$

The **limiting factor** is whichever timeline is shortest.




Figure 1: Effective constraint count $k_{\text{eff}} = k/(1 + \rho(k - 1))$ as a function of correlation ρ . **Key finding:** Regardless of how many raw constraints k a system accumulates, high correlation ($\rho \rightarrow 1$) crushes effective constraints to $k_{\text{eff}} = 1$. This is the “echo chamber” ceiling—adding more correlated constraints provides no marginal security. **Conclusion:** Scale alone cannot compensate for lost diversity.

4 Phase Space Analysis

4.1 The Chaos-Rigidity Spectrum

Systems can fail in two opposing directions:

Definition 4.1 (Chaos). A system is in a **chaos trajectory** when:

- Correlation $\rho < 0.2$ (constraints are independent)
- Entropy is increasing
- Defense function J exhibits high variance
- No emergent coordination

Definition 4.2 (Rigidity). A system is in a **rigidity trajectory** when:

- Correlation $\rho > 0.7$ (echo chamber formation)
- $k_{\text{eff}} \rightarrow 1$ regardless of k
- Single points of failure dominate
- Over-coordination suppresses diversity

In CCA, “rigidity” denotes an over-coordinated failure regime characterized by constraint redundancy and loss of adaptive capacity. Such systems fail due to Ashby-style requisite variety mismatch: homogeneous constraints cannot absorb heterogeneous perturbations.

Theorem 4.1 (Healthy Corridor). A system maintains coherence when:

$$0.2 < \rho < \rho_{\text{crit}} \quad \text{and} \quad \frac{\alpha}{k} \geq d \quad (18)$$

Remark 4.1 (Correction to v3). The second conjunct read $\alpha/k_{\text{eff}} > d$ in v3, inheriting the Theorem 2.3 error; it is corrected here to $\alpha/k \geq d$. The corridor bounds on ρ are unaffected. Note separately that the corridor edges are **substrate-specific** and require per-substrate calibration; the numerical bounds quoted here are GPU-anchored.




Figure 2: Phase space diagram showing the singularity boundary $K_{\text{req}} \cdot \rho = 1$ (purple curve). Above this curve, $T_{\text{truth}} \rightarrow \infty$ —the system cannot collapse incoherence regardless of scale or time. The healthy corridor exists between chaos ($\rho < 0.2$) and rigidity ($\rho > 0.7$), below the singularity.

4.2 Classification Algorithm

Algorithm: Trajectory Classification

Input: Timeline T_{eff} , measurements $\{(\rho_t, k_{\text{eff},t}, \sigma_t)\}$, thresholds $(\tau_{\text{collapse}}, \rho_{\text{low}}, \rho_{\text{high}}, \delta_\rho)$

Output: Classification $\in \{\text{HEALTHY}, \text{CHAOS}, \text{RIGIDITY}, \text{COLLAPSE}\}$

1. **If** $T_{\text{eff}} < \tau_{\text{collapse}}$ **then return** IMMINENT_COLLAPSE
2. $\rho_{\text{recent}} \leftarrow$ most recent ρ
3. $\text{trend}_\rho \leftarrow$ linear regression slope of ρ
4. **If** $\rho_{\text{recent}} < \rho_{\text{low}}$ **or** $\text{trend}_\rho < -\delta_\rho$ **then return** TRENDING_CHAOS
5. **If** $\rho_{\text{recent}} > \rho_{\text{high}}$ **or** $\text{trend}_\rho > +\delta_\rho$ **then return** TRENDING_RIGIDITY
6. **Return** HEALTHY

Remark 4.2 (Threshold Calibration). The algorithm requires four domain-specific thresholds:

- τ_{collapse} : Imminent collapse window (illustrative default: 7 days)
- ρ_{low} : Chaos boundary (illustrative default: 0.2)
- ρ_{high} : Rigidity boundary (illustrative default: 0.7)
- δ_ρ : Trend sensitivity (illustrative default: 0.01/day)

These are **not universal constants**. Calibration procedure:

1. **Historical fitting:** If collapse events exist in domain history, fit thresholds to maximize classification accuracy on labeled data.
2. **Sensitivity analysis:** Vary thresholds $\pm 20\%$ and assess classification stability.
3. **Domain expertise:** Adjust based on domain-specific knowledge (e.g., institutional systems may have higher ρ_{high} than ecological systems).
4. **Cross-validation:** Use held-out data to validate threshold choices.

The illustrative defaults are derived from sociotechnical systems with moderate coupling; other domains (biological, chemical, economic) may require substantial recalibration.

5 Information-Theoretic Limits

5.1 The L-01 Barrier

A fundamental result of CCA is that complete detection of incoherence is *provably impossible*.

Theorem 5.1 (L-01: Emergent Incoherence Barrier). There exists a non-empty class of emergent incoherence patterns that are fundamentally undetectable by any statistical analysis method operating on marginal distributions alone.

Proof. Incoherence can be *marginal-preserving*: each component’s output distribution matches the coherent distribution when viewed in isolation, while the joint distribution exhibits incoherence. Formally, let $P_{\text{honest}}(X_1, \dots, X_n)$ be the joint distribution of honest components, and let $P_{\text{deceptive}}$ be a deceptive joint distribution satisfying:

$$\forall i: P_{\text{deceptive}}(X_i) = P_{\text{honest}}(X_i) \quad (19)$$

Such distributions exist by construction (e.g., copula coupling with identical marginals). Any detector relying solely on marginal statistics has $I(X_i; \text{label}) = 0$, making detection impossible regardless of sample size. $\square$

Remark 5.1 (On the “40%” Estimate). The specific fraction of undetectable incoherence depends on the generative model for deception strategies. Under the assumption that detection operates via MI amplification with threshold ε , we model detection probability as $P(\text{detect}) = 1 - e^{-\beta\varepsilon}$ where β is a sensitivity parameter. For illustrative purposes, we use $\beta = 10$ and $\varepsilon \approx 0.092$, yielding $P \approx 0.60$. The complementary 40% represents the **residual undetectable fraction under this model**, not a universal constant. Different detection mechanisms or adversarial models would yield different fractions.

5.2 MI Amplification Detection

For non-marginal-preserving incoherence, detection is possible through mutual information (MI) amplification.

Proposition 5.2 (Detection Probability Model). Under an exponential detection model where the probability of detecting incoherence increases with MI amplification threshold ε , we propose:

$$P(\text{detect} \mid \varepsilon) = 1 - e^{-\beta\varepsilon} \quad (20)$$

where $\beta > 0$ is a sensitivity parameter that depends on the detector and signal characteristics.

Remark 5.2 (Model Justification). This functional form arises naturally from a Poisson detection model: if incoherence “events” arrive at rate proportional to ε , the probability of detecting at least one event in a fixed observation window is $1 - e^{-\beta\varepsilon}$. The parameter β must be calibrated empirically for each detection context.

Example 5.1 (Illustrative Calibration). With $\beta = 10$ (an illustrative choice):

- At $\varepsilon = 0.092$: $P(\text{detect}) = 1 - e^{-0.92} \approx 0.60$ (60%)
- At $\varepsilon = 0.050$: $P(\text{detect}) = 1 - e^{-0.50} \approx 0.39$ (39%)
- At $\varepsilon = 0.200$: $P(\text{detect}) = 1 - e^{-2.0} \approx 0.86$ (86%)

Proposition 5.3 (Multi-Layer Detection). With n independent detection layers, each with single-layer miss rate $q = 1 - P_1$:

$$P_{\text{combined}} = 1 - q^n \quad (21)$$

Layers	Detection Rate	Assumes $P_1 = 0.60$
1	60%	baseline
2	84%	$1 - 0.40^2$
3	94%	$1 - 0.40^3$
4	97%	$1 - 0.40^4$
5	99%	$1 - 0.40^5$

Table 3: Multi-layer detection improves coverage but cannot reach 100%. Rates assume 60% single-layer detection; actual rates depend on calibrated P_1 .

6 Intervention Analysis

CCA identifies intervention types and their expected effects on the defense function.

Intervention	Target	Effect	Cost
Add diverse agents	$\rho \downarrow, k \uparrow$	$k_{\text{eff}} \uparrow$	High
Increase strictness	$\lambda \uparrow$	$J \uparrow$	Medium
Generate sustainability	$\sigma \uparrow$	$T_{\text{entropy}} \uparrow$	Low
Rotate nodes	$f \downarrow$	$T_{\text{capture}} \uparrow$	High
Prune constraints	$k \downarrow$	Reset α/k	Medium

Table 4: Intervention types and their effects.

6.1 Intervention Selection

Definition 6.1 (Priority Score). The **priority score** for an intervention is:

$$\text{Priority} = \frac{\Delta J}{\text{Cost}} \quad (22)$$

The optimal intervention depends on the current trajectory:

- **Trending chaos**: Increase strictness, add structure
- **Trending rigidity**: Add diverse agents, reduce correlation
- **Imminent collapse**: Emergency intervention on limiting factor

6.2 A Practical Intervention: CIRISAgent

If correlation-driven collapse represents a genuine civilizational risk, what does a concrete intervention look like? The companion paper “*CIRISAgent: An Open-Source Framework for Ethical AI Through Transparent Architecture*” presents one implementation—an AI system designed explicitly to resist the rigidity trajectory.

CCA Intervention	CIRISAgent Component		Mechanism
Add diverse agents ($\rho \downarrow$)	22-service architecture	microarchi-	Modular services prevent monolithic correlation
Increase strictness ($\lambda \uparrow$)	H3ERE Module	Conscience	Coherence faculty detects contradictions
Generate sustainability ($\sigma \uparrow$)	Gratitude token system		Positive-sum incentives vs zero-sum capture
Rotate nodes ($f \downarrow$)	Wise Authority Service		Human oversight prevents Byzantine capture
Transparency (detection $\uparrow$)	Audit & Visibility Services		Hash-chained logs expose emergent deception

Table 5: CIRISAgent components mapped to CCA intervention strategies.

Key design principles for collapse resistance:

1. **Transparency over opacity:** Every decision is auditable. Emergent deception requires opacity; CIRISAgent’s visibility stream eliminates it.
2. **Pluralistic oversight:** The Wise Authority Service mandates human escalation for high-stakes decisions, preventing single-point-of-failure rigidity.
3. **Modular independence:** The 22-service architecture ensures that constraint correlation (ρ) remains low—services can be updated, replaced, or audited independently.
4. **Explicit deferral:** Agents recognize limitations and defer rather than optimize past competence boundaries.

Remark 6.1 (AI as Correlation Amplifier). The CCA framework identifies AI systems as significant not because they are uniquely dangerous, but because they **accelerate correlation faster than any prior technology** [41, 43]. Empirical evidence shows LLM outputs converge toward “bland central tendencies” [40], human-AI loops amplify bias beyond human-human baselines [39], and algorithmic monoculture creates systemic risk [45]. A naive AI optimization loop drives $\rho \rightarrow 1$ while appearing to increase k —the appearance of diversity with the reality of monoculture. CIRISAgent is designed to break this pattern through architectural constraints that preserve effective diversity (k_{eff}) even as nominal capability (k) increases.

CIRISAgent does not claim to solve alignment. It implements one trajectory through the design space that prioritizes collapse resistance over raw capability. The framework is open-source, auditable, and explicitly incomplete—an invitation to extend rather than a claim of sufficiency.

Resources:

- CIRISAgent implementation: <https://github.com/CIRISAI/CIRISAgent>
- RATCHET validation framework: <https://github.com/CIRISAI/RATCHET>
- Research status: <https://ciris.ai/research-status/>

7 Validation

7.1 Monte Carlo Self-Consistency Check

Version 3 reported this table as Monte Carlo *validation* of the detection probability model. It is not a validation, and the label is corrected here. The simulation draws detection events with

probability $1 - e^{-\beta\hat{\varepsilon}}$ from a noisy estimate $\hat{\varepsilon}$ and compares the resulting rate against $1 - e^{-\beta\varepsilon}$ — the same expression at the noiseless value. The two cannot disagree except through the concavity of $1 - e^{-\beta\cdot}$ in its argument, so the residual column measures a Jensen bias, not model adequacy. The closed-form prediction for that bias, $-\frac{1}{2}\beta^2 e^{-\beta\varepsilon}\sigma^2$, reproduces the observed differences: -1.21% , -0.80% , -0.45% , -0.27% against the tabulated -0.4% , -0.9% , -0.5% , -0.4% . **No independent detection process is involved, and the reported RMSE of 0.006 measures how faithfully the sampler reproduces its own sampling distribution.** This is the third instance in this paper of a residual that is a Jensen gap rather than a measurement error (cf. Remark 10.4).

What the table legitimately shows is that the implementation samples from the model it claims to sample from, with $\beta = 10$ and measurement noise $\sigma = 0.02$.

ε	Theoretical $1 - e^{-10\varepsilon}$	Empirical	Difference
0.050	39.3%	38.9%	−0.4%
0.092	60.1%	59.2%	−0.9%
0.150	77.7%	77.2%	−0.5%
0.200	86.5%	86.1%	−0.4%

Table 6: Monte Carlo self-consistency check, $\beta = 10$. **Not a validation:** the “empirical” column is sampled from the same expression the “theoretical” column evaluates. The difference column is the concavity bias of the noise in $\hat{\varepsilon}$, matching the closed-form Jensen prediction, and the RMSE of 0.006 measures sampler fidelity rather than model adequacy.

7.2 Formal Verification

Key theorems have been formalized in Lean 4 with Mathlib dependencies:

- `EmergentDeceptionBounds.lean`: L-01 impossibility theorem
- `MIAmplificationTheory.lean`: Detection probability bounds
- `Chronometry.lean`: Collapse timeline definitions

Representative theorem statement (singularity condition):

```
theorem CT1_singularity_boundary
  (K_req rho : R) (h_pos : 0 < K_req) (h_bound : K_req * rho >= 1) :
  T_truth K_req rho alpha = top := by
  -- When denominator (1 - K_req * rho) <= 0, T_truth is undefined/infinite
  unfold T_truth
  simp [h_bound, le_antisymm]
```

Note: The singularity arises because the denominator $(1 - K_{\text{req}} \cdot \rho)$ becomes zero or negative when $K_{\text{req}} \cdot \rho \geq 1$. In the formalization, this is represented as $\top$ (infinity in the extended reals), indicating that the system *cannot* collapse deception in finite time.

Remark 7.1 (Scope of Formal Verification). Formal verification in Lean 4 proves **internal consistency**: *if* the world behaves according to the CCA equations, *then* the derived theorems hold. It does **not** prove external validity—that is, whether the CCA equations accurately model real-world systems. The proofs establish that our mathematics is self-consistent, not that our assumptions are correct. Empirical validation (cross-domain testing, Monte Carlo simulation) addresses external validity; formal verification addresses logical soundness.

The complete implementation, including Python modules, Lean proofs, and simulation scripts, is available at github.com/CIRISAI/RATCHET under AGPL-3.0 license.

8 Cross-Domain Validation

A critical test of CCA’s claims is whether the framework generalizes beyond sociotechnical systems. If the k_{eff} formula and collapse dynamics reflect universal properties of constraint-based systems, they should apply across fundamentally different scientific domains. We implemented three domain-specific simulation engines and validated them against authoritative empirical datasets.

8.1 Domain Mapping

Each engine maps domain-specific variables to CCA structural invariants:

CCA Variable	Battery	Institutional	Microbiome
k (constraints)	Cell count	Executive constraints	Species count
ρ (correlation)	Cross-cell SOH correlation	Elite coupling	SparCC correlation
σ (sustainability)	State of Health	Political stability	Shannon diversity
f (compromise)	Capacity fade	Corruption fraction	Pathogen fraction
α (generation)	SEI growth rate	Reform rate	Colonization rate
d (decay)	Calendar aging	Institutional erosion	Substrate decay
λ (strictness)	Operating window	Rule of law	Interaction strength

Table 7: CCA variable mappings across chemistry, political science, and biology domains.

8.2 Empirical Data Sources

Three authoritative datasets provide ground truth:

- **NASA Li-ion Battery Aging Dataset:** 19 cells with charge/discharge cycling and impedance measurements (NASA Prognostics Center of Excellence).
- **Quality of Government + Polity V:** 203 countries, 12,393 observations (1946–2024), with 398 regime collapse events coded from Polity transition indicators.
- **American Gut Project:** 100 samples, 2,081 microbial taxa with genus-level taxonomic resolution.

8.3 k_{eff} Formula Application

The formula $k_{\text{eff}} = k/(1 + \rho(k - 1))$ takes a constraint *count* $k \geq 1$ and an intraclass correlation ρ among those k constraints. Two of the four domain rows carried in v3 supplied a genuine count; the two institutional rows did not and are withdrawn (Remark 8.1).

Domain	k	ρ	k_{eff}	Interpretation
Battery (NASA, 19 cells)	19	0.000	19.00	Cells in the pack (integer count)
Microbiome (AGP average)	365.7	0.190	5.20	Observed species per sample, averaged

Table 8: k_{eff} applied in the two domains supplying a constraint count. **Two caveats.** (i) k_{eff} is *computed* from k and ρ and is not independently observable, so this table exhibits the formula’s behaviour rather than testing it. (ii) Neither ρ is an independently estimated intraclass correlation: the battery figure is a dispersion transform of final-SOH spread across cells, and the microbiome figure is a deterministic function of the same Shannon diversity reported as σ , confined by construction to $[0.14, 0.30]$. Both are stated here as proxies. (iii) **The microbiome row nests one effective-count transform inside another.** Its σ is Shannon diversity, whose own standard effective-number conversion is $\exp(H)$ [12]; feeding it to k_{eff} as a raw input therefore composes two effective-count constructions without saying so, and the resulting $k_{\text{eff}} = 5.20$ should not be read as an independent effective species count. Read this table as a provenance record, not as cross-domain validation.

Key finding: The formula reduces the effective constraint count as correlation rises, identically in both domains where k is a count. When $\rho \rightarrow 0$, $k_{\text{eff}} \rightarrow k$; when $\rho \rightarrow 1$, $k_{\text{eff}} \rightarrow 1$. This is an algebraic property of the Kish design effect, *exhibited* here rather than established here.

Remark 8.1 (Correction to v3: the institutional rows are withdrawn, and not recomputed). Version 3 carried two institutional rows: Venezuela at $k = 0.667$, $\rho = 0.299$, $k_{\text{eff}} = 0.55$, and Turkey at $k = 1.000$, $\rho = 0.000$, $k_{\text{eff}} = 1.00$. Both are withdrawn.

The k values are not counts. They are the Polity executive-constraints ordinal rescaled to the unit interval, $k = (\text{xconst} - 1)/6$: Venezuela’s 0.667 is $\text{xconst} = 5$, Turkey’s 1.000 is $\text{xconst} = 7$. Kish’s k is a cluster count with $k \geq 1$, and outside that domain the formula does not merely lose accuracy — it *reverses*. For $k < 1$, $\partial k_{\text{eff}}/\partial \rho = -k(k-1)/(1+\rho(k-1))^2 > 0$, so k_{eff} *increases* with correlation, rising from 0.667 at $\rho = 0$ to 1.000 at $\rho = 1$. Venezuela’s stated reading, “elite coupling reduces diversity,” therefore asserts the opposite of what the formula does at those inputs. The printed 0.55 lies below that entire attainable range, so no value of ρ produces it; correct arithmetic on the printed inputs gives 0.7408, which would be a correct calculation on inapplicable numbers.

The ρ values are not correlations. They were computed as $\rho = f(1-k)$ — a corruption index times a constraint deficit, with no covariance entering. Turkey’s $\rho = 0.000$ is thus *forced*: at $\text{xconst} = 7$ the deficit vanishes and ρ is identically zero whatever the corruption level. Six of the fourteen countries in the institutional analysis are zeroed this way. That row is doubly uninformative, since the Kish map is also constant in ρ at $k = 1$.

Why these rows are not reconstructed. The obstacle is not merely that the sources used here supply no count — though they do not; the only count-valued institutional variable available is the number of legislative chambers, which is 1 for both countries throughout the window and returns the same degenerate case. The deeper obstacle is that the design effect assumes the k units are *exchangeable*: draws from a common pool sharing one pairwise correlation, which is what makes ρ an *intraclass* correlation. Executive, legislature and judiciary are structurally different kinds of institution, not exchangeable draws. The equicorrelated precondition therefore fails by construction for this object, independently of data availability — and it is the same precondition that §10.2 finds already strained on the GPU substrate at low ρ . Supplying a veto-player count from another source would not repair this; it would reintroduce the error with better inputs. See CORRECTIONS_v3.md (C-2, C-3, C-15).

8.4 Empirical Performance

CCA excels at **structural diagnosis** (identifying phase and trajectory) and shows higher uncertainty in **timing estimates**. This aligns with its design as a risk framework: it answers “is this system fragile?” more reliably than “when will it fail?”

8.4.1 Battery Engine vs NASA Data

- **SOH RMSE**: 8.1% average across 19 cells
- **σ validation**: Computed $\sigma = 79.31\%$, matches manual calculation exactly
- **WITHDRAWN**. v3 reported “ f validation: computed $f = 20.69\% = 1 - \sigma$ (exact match).” In the engine f is *defined* as $1 - \sigma$ (`battery.py:427` returns `1 - self.get_sigma()`; the neural engine assigns `f = max(0, 1.0 - sigma)`), so the “exact match” reports that $1 - \sigma = 1 - \sigma$. This is a definition, not a validation, and no measurement could have contradicted it
- **Limitation**: Model overestimates final SOH by $\sim 8\%$ due to simplified SEI kinetics

8.4.2 Institutional Engine vs QoG/Polity

Withdrawn from this chapter. Version 3 reported here that the institutional engine correctly classified 5/5 stable democracies and detected fragility a mean 7.6 years early. Re-scoring against a pre-specified outcome definition does not support either claim, and the engine does not consult k_{eff} at all, so its performance — whatever it were — would not bear on the framework. The case study is retained, honestly scored, as a negative result in §11, outside the validation chapter.

8.4.3 Microbiome Engine vs American Gut Project

- **k range**: 238–643 observed species (matches AGP literature)
- **σ mean**: 0.580 Shannon diversity (normalized), consistent with healthy adult norms
- **f mean**: 0.232 pathogen fraction
- **ρ mean**: 0.19 (moderate community coupling)

8.5 Testing Rigor: Bug Discovery

Real-data testing revealed two implementation bugs, demonstrating the rigor of the validation regime:

1. **Antibiotic shock did not reduce k** : Species abundances were reduced but not eliminated. *Fix*: Added extinction threshold check after shock application.
2. **FMT intervention decreased σ** : Default donor profile used high-variance lognormal distribution (low diversity). *Fix*: Changed to low-variance lognormal (high diversity, matching healthy donors).

These were *logic bugs*, not parameter calibration issues, and were corrected without adjusting any parameters that would constitute gaming results.

8.6 Structural Invariant Summary

Remark 8.2 (Correction to v3: this subsection is withdrawn). Version 3 reported that “three CCA invariants were validated across all domains,” listing I-01 (k_{eff} identity, “zero numerical error across chemistry, biology, and political science”), I-02 ($f = 1 - \sigma$), and I-03 (a collapse threshold). **No validation was performed.** The cited artifact emits the three findings as fixed text; it reads no data, compares nothing, and has no execution path that could report anything

else. I-02 is additionally an assignment statement in the source ($f := \max(0, 1 - \sigma)$), so its “exact match” is the assignment operator. I-01 restates the identity discussed at Remark 10.1.

The subsection is withdrawn in full. Nothing in it constitutes evidence, and the cross-domain generality it asserted is not established by this paper. See `CORRECTIONS_v3.md` (C-12).

Remark 8.3 (Correction to v3: the cross-domain generality claim is withdrawn). Version 3 closed this chapter by reading the three domain applications as evidence that CCA captures structural properties common to constraint-based systems rather than domain-specific phenomena, concluding that “three domains demonstrate cross-domain applicability.” That conclusion is withdrawn on three independent grounds. The institutional application is withdrawn from this chapter entirely and reported as a negative result (§11). The structural-invariant subsection that supported it computed nothing. And applying one formula to three datasets demonstrates that the formula *can be applied*, not that it *holds* — a distinction the remainder of this chapter now states directly. What survives is two domain applications, both carrying proxy ρ values rather than measured intraclass correlations.

9 Related Work

CCA synthesizes four research traditions while contributing novel formal structures not present in prior work.

9.1 Safety Engineering: FMEA and Fault Tree Analysis

Failure Mode and Effects Analysis (FMEA), developed by the U.S. military in the 1940s, provides systematic enumeration of failure modes with severity/likelihood scoring. Recent applications extend to cybersecurity governance under NIS 2 and DORA regulations, and to enterprise risk management in financial contexts. CCA adopts FMEA’s structured failure enumeration but contributes *closed-form collapse timelines* with identified singularities—a quantitative precision absent from traditional FMEA worksheets.

9.2 Control Theory and Stability Analysis

The stability condition $\alpha/k \geq d$ (Theorem 2.3) follows the control-theoretic tradition of identifying regions where system dynamics remain bounded. Stafford Beer’s Viable System Model (VSM), building on Ashby’s Law of Requisite Variety, established that “only variety can absorb variety”—a controller must match the complexity of its environment. CCA formalizes this insight: the rigidity boundary ($k_{\text{eff}} \rightarrow 1$) represents a formal instance of loss of requisite variety, where correlated constraints cannot absorb environmental complexity regardless of scale. The “complexity misalignment” literature identifies this failure mode conceptually; CCA provides the singularity condition $K_{\text{req}} \cdot \rho \geq 1$ as its formal expression.

9.3 Network Collapse and Phase Transitions

Recent work on explosive synchronization demonstrates that collapse and recovery trajectories depend on proximity to first-order phase transitions. A 2025 PNAS study shows this framework predicts consciousness loss under anesthesia and market collapse during the 2008 crisis. The Watts cascade model and Granovetter threshold models explain how small shocks trigger large cascades in social systems. CCA extends this tradition by identifying *three independent collapse modes* ($T_{\text{truth}}, T_{\text{entropy}}, T_{\text{capture}}$) with distinct dynamics, rather than a single cascade mechanism. Importantly, these modes can dominate at different timescales, a feature not captured by single-mechanism cascade models. The chaos–healthy–rigidity phase space provides explicit boundaries absent from prior cascade models.

9.4 Entropy and Decoherence in Organizations

The nearest conceptual neighbor is recent work on “Decoherence in Socio-Technical Organisations” (December 2025 preprint), which models coordination breakdown using entropy and cognitive burden concepts. This work shares CCA’s concern with coherence loss but does not provide the three-clock framework, singularity conditions, or information-theoretic detection bounds that distinguish CCA. The sustainability transitions literature similarly addresses sociotechnical system destabilization but remains largely descriptive rather than formally predictive.

9.5 Byzantine Fault Tolerance and Capture Dynamics

The T_{capture} timeline builds on Byzantine fault tolerance (BFT) foundations, where $n \geq 3f + 1$ nodes are required to tolerate f Byzantine failures. Recent work addresses node capture attacks through reputation-enhanced PBFT (RePA, 2025) and dynamic consensus algorithms (LTSBFT, 2024). However, these focus on *detection and mitigation* of captured nodes rather than *prediction of capture timelines*. CCA contributes the explicit formula $T_{\text{capture}} = (n/3 - f)/r_{\text{cap}}$, treating capture as a race condition between adversarial progress and system response.

9.6 Information-Theoretic Detection Limits

Federated learning security research reports detection accuracies of 85–89% under normal conditions and 66–73% under adversarial attack, focusing on empirical detection performance rather than fundamental limits. L-01 sits in the latter category, but its scope must be stated precisely: it establishes that the class of marginal-preserving incoherence patterns undetectable from marginals alone is **non-empty** — an existence result, with **no measure**. The “~40%” figure in earlier versions of this paper followed from one illustrative parameterization ($\beta = 10$, $\varepsilon \approx 0.092$) and is a stated wager, not a derived bound; the detection model $P(\text{detect}) = 1 - e^{-\beta\varepsilon}$ is a modelling choice with a free sensitivity parameter, not a measurement.

Adjacent formal work is stronger on this point and should be read alongside it: Dorner et al. [16] establish a ceiling on LLM-judge evaluation — when the judge is no more capable than the model being evaluated, no debiasing procedure halves the required ground-truth data. That is a quantified limit; L-01 is an existence claim. The two should not be cited as comparable.

9.7 Lineage and Contribution

Earlier versions of this section claimed, on the basis of a search of 2024–2025 literature, that no prior work combined this paper’s elements. A two-year window is not a prior-art search, and the claim is withdrawn. The central insight here is old, well established, and belongs to several literatures at once.

That correlated voters degrade collective accuracy is settled, and has been since 1992. Ladha [2] extended the Condorcet Jury Theorem to correlated votes; Ladha [4] generalized it across normal, hypergeometric and Pólya correlated-voting models and proved that the effectiveness of majority-rule voting *decreases* as correlation increases, further arguing that diversity among decision-makers improves decisions by lowering average correlation. Austen-Smith and Banks [5] sharpen this: independence is *necessary*, not merely sufficient, and correlated voters can perform *worse* than individuals — a stronger statement than the singularity condition derived in this paper, and one that should be read as bounding it.

That nominally independent implementations fail together is settled empirically, and has been since 1985. Eckhardt and Lee [6] gave the theoretical basis; Knight and Leveson [7] wrote 27 versions independently from one specification, ran a million tests, and found coincident failures substantially exceeding the independence prediction, defending the result against criticism in a later reply [8]. The finding has recently been replicated on LLM

coding agents [9], where 429 coincident-failure cases arose against an independence prediction of 115.36.

The effective-count-under-correlation object itself is standard across fields. Kish [1] in survey sampling; Laakso and Taagepera [10] as the effective number of parties (equivalent to the inverse Simpson and inverse Herfindahl–Hirschman indices); Hill [11] and Jost [12] as Hill numbers in ecology; Meucci [13] as the effective number of bets in portfolio theory. These are the same construction reached independently in four disciplines.

In the AI setting the ground is likewise occupied. Lefort et al. [14] applied the Condorcet Jury Theorem to LLM ensembles and found marginal gains implying a lack of independence; Kim et al. [15] measured error correlation across 350+ models; Jiang et al. [17] documented output homogeneity within and across model families. Concurrently and independently, Kohli [18] applies the Kish identity to a 9-judge, 7-family LLM panel, reporting $n_{\text{eff}} = 2.18$ at mean pairwise $\phi = 0.391$ with a hard asymptote at $1/\bar{\phi} \approx 2.6$ — the Möbius ceiling of §2.2, independently obtained. That work does not cite this one, and is not derivative of it.

What remains this paper’s own is narrower than earlier versions claimed, and is stated here at its true size: the three closed-form collapse timescales, the singularity boundary and the phase corridor built on that base, and a production measurement substrate. The underlying correlation result is inherited, not introduced, and is cited above as such.

10 Physical Validation: GPU Timing Strain Gauge

The theoretical framework of CCA was validated physically using a GPU timing strain gauge—a 128-sensor array measuring kernel timing jitter across an NVIDIA RTX 4090 die. This provides the first *hardware instantiation* of CCA dynamics, demonstrating that correlation-driven diversity collapse is not merely a mathematical abstraction but a measurable physical phenomenon.

10.1 Experimental Setup

Each “sensor” measures nanosecond-resolution kernel timing. The array computes:

- k : Number of sensors (128)
- ρ : Average pairwise timing correlation
- k_{eff} : Computed via the Kish formula

10.2 Kish Identity: Implementation Check (Exp 86)

Version 3 reported the core formula as validated with “perfect correlation.” That report was wrong in kind; the table is retained below with its correction:

Metric	Predicted	Measured	Correlation
$k_{\text{eff}} = k/(1 + \rho(k - 1))$	–	–	$r = 1.000$
Baseline ρ	–	0.13	–
Baseline k_{eff}	7.5	7.5	exact

Table 9: Exp 86 Kish comparison. **Withdrawn as validation** — see Remark 10.1.

Remark 10.1 (Correction to v3: this is an implementation check, not a validation). In the Exp 86 instrument, the quantity labelled *measured* k_{eff} and the quantity labelled *theoretical* k_{eff} are computed by the same expression $k/(1 + \rho(k - 1))$ from the same ρ . The reported figure is $\text{corr}(f(\rho), f(\rho))$, which returns $r = 1.000$ for *any* formula whatsoever and therefore carries no

evidence about this one. Table 8 states the same point directly: k_{eff} is computed from measured k and ρ and is not independently observable.

What Exp 86 does establish, and no more: the implementation computes the identity it claims to compute. That is a software correctness check, belonging with the Lean formalization of §7.2 rather than with empirical results.

The one falsifiable prediction in Exp 86 failed, and v3 did not report it. Exp 86 predicted that common-mode injection at strength s drives the array to $\rho = \rho_0 + s(1 - \rho_0)$. Measured against that prediction the injection model achieves $R^2 = 0.714$ with maximum absolute error 0.234 in ρ , overstating the approach to full correlation by up to 47% at mid-range. Version 3 reported the identity that could not fail and omitted the model that did.

The baseline figures require two further corrections. They originate in Exp 84 (a 128-sensor array), not Exp 86 (64 sensors), whose own baseline is $\rho = 0.167$, $k_{\text{eff}} = 5.55$. And Exp 86 estimates ρ as the mean *absolute* pairwise correlation, which is not the signed intraclass correlation the Kish identity is defined over and which carries a positive finite-sample floor: at 64 sensors and 50 samples that estimator reads ≈ 0.114 on synthetic data with true $\rho = 0$, implying $k_{\text{eff}} \approx 7.8$ from independence alone. The reported baseline is not separated from that floor, and no shuffle null was computed. See `CORRECTIONS_v3.md` (C-5).

Remark 10.2 (What a genuine validation of the Kish relation would require). The identity becomes empirically testable only against an *independently observed* measure of effective diversity — one not computed from ρ . Candidates: the variance of the array mean ($\sigma_{\bar{x}}^2 = \sigma^2/k_{\text{eff}}$), the SNR gain from averaging N sensors, the participation ratio of the correlation spectrum, or a task-level measure such as detection latency. **This paper reports no such comparison**, and neither Exp 86 nor C1 supplies one.

One of those candidates does exist on other data, and is cited here because it is the nearest thing to the missing measurement. Moore [22] computes effective dimensionality on $n = 6,465$ production reasoning traces from a deployed CIRIS agent using the **participation ratio** of the eigenvalue spectrum of a standardized feature covariance, $N_{\text{eff}}^{\text{PR}} = (\sum_i \lambda_i)^2 / \sum_i \lambda_i^2$, alongside an entropy-perplexity variant. This is computed from the covariance structure directly and **not from a scalar ρ** , so — unlike Exp 86 and C1 — it is free to disagree with a Kish value. It reports a 99% variance horizon of 11 dimensions, a 90% horizon of 7, a mean participation ratio of 6.61, and an outcome link: traces at $N_{\text{eff}} \geq 9.2$ overrode restrictive foundation-model priors at 83% reliability, benchmarked against response-entropy and latency baselines.

Cited with its limits, which are material. The reported stability threshold ($N_{\text{eff}} \approx 7.1$) did not come from a held-out cohort; that paper states it “emerged naturally” through continuous iterative tuning of the decision-making algorithms and conscience module, so 7.1, 9.2 and the 83% are in-sample. The threshold is also stated across two estimators: the mean participation ratio is 6.61, *below* 7.1, while the “ ≥ 8.5 consistently maintained” figure is the entropy-perplexity estimator rather than the participation ratio, and which of the two the threshold is defined over needs resolving. Finally, the 90% variance horizon (7 dimensions) sits within rounding of the threshold (7.1), so it is not yet ruled out that the threshold recovers the intrinsic rank of the reasoning manifold rather than marking a dynamical boundary — which would make the claim a restatement of the manifold’s dimensionality rather than a finding about deception.

What this closes and what it does not. It supplies a spectral effective-dimensionality measurement, with an outcome link, on production data — materially more than this paper has. It does **not** validate the Kish identity: it is a different functional on a different substrate, and no paired comparison of $N_{\text{eff}}^{\text{PR}}$ against k_{eff} on the same data has been run. The GPU array is where that comparison is cheapest, since both quantities can be computed from correlation matrices already stored, and it remains the outstanding work (§12.2).

The nearest available evidence bears on the identity’s **precondition** rather than the identity. The Kish form assumes an equicorrelated correlation matrix. Exp 117 tests that assumption by comparing the spectral entropy deficit computed from all 16 eigenvalues of the measured matrix

against the closed form predicted from the scalar ρ alone — distinct functionals over distinct inputs, free to disagree. They agree to 0.8% at $\rho = 0.394$, and diverge by 17.6% and 70.8% at $\rho \approx 0.04$. The equicorrelated assumption is therefore **supported at high correlation and falsified at low correlation**, in the regime where the estimator floor above dominates. The failure is reported alongside the success: it bounds the domain over which any k_{eff} figure in this paper should be believed.

10.3 Software-Induced Collapse (Exp 103) — Withdrawn

Version 3 reported that synchronized GPU workloads demonstrate the singularity boundary ($\rho \rightarrow 1 \Rightarrow k_{\text{eff}} \rightarrow 1$), collapsing the array to $k_{\text{eff}} = 1$ through software coordination alone. **That result is withdrawn.** The v3 table is reproduced here so the correction is checkable against what it corrects:

Workload (v3)	ρ	k_{eff}	Status now
Baseline (normal)	0.14	6.5	Contaminated by warm-up (below)
Barrier sync	0.90	1.1	Single draw; not reproducible as a point value
Lockstep kernels	1.00	1.0	Measurement artifact

Table 10: The v3 Exp 103 table, retained for reference. None of the three rows survives as reported.

Remark 10.3 (Correction to v3: replication on the original hardware). The experiment was re-run on the same GPU under a pre-registered protocol (predictions, interpretations and a decision rule fixed before any measurement; the registration is hash-pinned and cited in `CORRECTIONS_v3.md`). Three findings.

The lockstep row is an artifact of the measurement code. The lockstep routine times one batch and assigns that single scalar to all 64 sensors, so the correlation matrix is computed over identical rows. Replication confirms the rows are byte-identical, with $\rho = 1.0000$ exactly and zero variance after common-mode removal. $k_{\text{eff}} = 1$ here is arithmetic on copied numbers, not a collapse. This is the same defect class as Remark 10.1 and Remark 10.4.

The barrier effect is real but carries no point estimate. Barrier synchronization does raise correlation: $\rho_{|r|} = 0.79$ against a shuffle-null 95th percentile of 0.06, and it survives common-mode removal at 0.44 — so this is genuine structure, not estimator floor, and v3’s qualitative direction is supported. But across five identical trials at fixed parameters the same procedure returned $\rho \in [0.278, 0.809]$, overlapping the independent condition’s own range $[0.111, 0.777]$. In k_{eff} terms that spans 1.23 to 3.46: whether a given run reads as “collapse” depends on which draw it got. The v3 value of 0.90 is one draw from that distribution, reported without repetition or interval.

The baseline is contaminated by warm-up. The first measurement taken after the GPU has been idle reads $\rho \approx 0.78$ regardless of condition, settling to ≈ 0.11 – 0.17 on subsequent runs; this was observed in two independent sessions. Exp 103 measures its baseline first, so the 0.14 baseline and every $\Delta\rho$ computed against it are unreliable.

Consequence. The claim that *software coordination alone collapses effective diversity* is not supported: its demonstration was the lockstep artifact, and the barrier effect — though real — cannot carry a collapse claim at this variance. No prior run in this series computed a null, repeated a measurement, or reported an interval; at the shipped sample count the shuffle null’s 95th percentile is 0.15, so some portion of the correlations reported across these experiments may be estimator floor. See `CORRECTIONS_v3.md` (C-13).

10.4 Recovery Dynamics (Exp 89)

Post-collapse recovery was measured at $\tau = 6.5$ ms (electrical timescale), confirming that the collapse is reversible when the correlation-inducing workload is removed. This validates CCA’s assumption that correlation is a dynamic quantity, not a fixed system property.

10.5 Leading Indicator Detection (Exp 88)

Early warning signals appeared at $\rho = 0.28$ —well below the $\rho_{\text{crit}} = 0.43$ collapse threshold—providing 7 measurement cycles of advance warning. This validates CCA’s classification algorithm: trend detection ($\Delta\rho > 0$) provides actionable lead time before threshold crossing.

10.6 Signal Architecture

The GPU experiments revealed a dual-output architecture that mirrors CCA’s dual concern with entropy and correlation:

Output	Source	CCA Mapping
TRNG (entropy)	Raw timing LSBs (99.5% white noise)	System entropy σ
Strain gauge	k_{eff} dynamics (reset cycles)	Correlation tracking ρ

Table 11: GPU timing provides both entropy generation and correlation measurement from a single physical source.

Critical finding: Environmental signals (thermal $r = 0.30$, EMI $r = 0.21$) are undetectable in raw timing but visible through oscillator dynamics. The oscillator acts as a *correlation detector*, converting amplitude modulation (buried in noise) to timing modulation (measurable).

10.7 Mean Shift Detection (O-Series Validation, January 2026)

Cross-team validation (Array B1e $\rightarrow$ Ossicle O1-O5) revealed that **timing mean shift**, not variance ratio, is the optimal workload detection signal:

Intensity	Mean Shift	Regime	Detection
1%	+191%	Binary	YES
10%	+198%	Binary	YES
30%	+215%	Transition	YES
50%	+248%	Linear	YES
90%	+380%	Linear	YES

Table 12: Two-regime detection model (O5). Workloads as low as 1% cause +191% mean shift due to scheduler/arbitration overhead.

Key findings:

- **Detection floor:** 1% workload (not 30% as previously claimed)
- **Two-regime model:** Binary threshold ($< 30\%$: flat $\sim 200\%$ shift) + linear scaling ($> 30\%$: $\text{shift} = 227 \times \text{intensity} + 160$)
- **Optimal sample rate:** 4000 Hz (lowest variance $\pm 14\%$); avoid 1900–2100 Hz (interference dip)
- **Detection threshold:** $> 50\%$ mean shift (provides margin for background tasks)

The two-regime behavior reflects GPU contention physics: *any* workload triggers scheduler arbitration overhead (+160% baseline), with linear scaling only above 30% intensity.

10.8 Coherence Collapse Propagation (C-Series, January 2026)

The C-series experiments used a 16-sensor 4×4 array at 9631 Hz to measure the *spatial dynamics* of coherence collapse—the first physical characterization of how correlation propagates across a semiconductor die.

10.8.1 C1: k_{eff} Formula Empirical Validation

The Kish formula was tested by inducing correlation through barrier synchronization across 21 `sync_strength` levels from 0.0 to 0.8:

ρ	Measured k_{eff}	Predicted k_{eff}	Δ
0.02	12.4	11.8	+0.6
0.10	6.8	6.4	+0.4
0.20	4.6	4.0	+0.6
0.30	3.2	2.9	+0.3
0.40	2.5	2.2	+0.3
<i>(representative subset; full data: $n = 21$ points)</i>			

Table 13: C1 k_{eff} comparison across 21 induced correlation levels ($\rho \in [0.02, 0.40]$). **Withdrawn as validation** — see Remark 10.4.

Remark 10.4 (Correction to v3: C1 is an identity check, not a validation). Version 3 reported this comparison as empirical validation of the k_{eff} formula on physical hardware ($R^2 = 0.798$, $n = 21$). **It is not a test of the formula.** Both columns are generated by $k/(1 + \rho(k - 1))$ from the same measured correlation matrix, differing only in the order of averaging: the column labelled *measured* is $\frac{1}{k} \sum_i f(\rho_i)$ (the identity applied per sensor, then averaged) and the column labelled *predicted* is $f(\bar{\rho})$ (the identity applied to the average). Because f is convex in ρ , Jensen’s inequality guarantees the first exceeds the second at every point — and it does, in **21 of 21 trials**, mean +0.76, minimum +0.30, never once crossing zero. A genuine measurement-versus-prediction comparison scatters about zero. The residual reported as $R^2 = 0.798$ is the magnitude of a convexity gap: a property of the arithmetic, not of the hardware.

C1 is retained as a measurement of *induced correlation versus barrier-sync strength*, which is a real result. Its k_{eff} column is a derived restatement of its ρ column and is withdrawn as independent evidence. See `CORRECTIONS_v3.md` (C-11).

10.8.2 C2: Correlation Propagation Velocity

Collapse was induced at one corner of the array and propagation to distant sensors was timed:

Metric	Value
Propagation velocity	0.5 ± 0.4 m/s
Die crossing time	36.9 ms
Regime	Thermal (not electrical)
Measurements	120 (10 trials $\times$ 12 sensors)

Table 14: Correlation propagation velocity (C2). First measurement on GPU die.

Key finding: Correlation propagates at ~ 0.5 m/s—intermediate between thermal diffusion (~ 0.01 m/s) and electrical speeds ($\sim 0.1c$). This means coherence collapse unfolds over tens of milliseconds, providing a physically-grounded timescale for intervention.

Note on mechanism (E3 follow-up): The measured 0.5 m/s is $\sim 50\times$ faster than expected from thermal diffusion alone. E3 experiments found: (1) sample rate has no effect (ruling out sensor latency), (2) thermal vs algorithmic stimulus showed no difference, (3) spatial variation $CV=0.54$ suggests non-uniform propagation. *One hypothesis* is that correlation propagates through shared GPU resources (L2 cache, memory controller, power delivery network) rather than physical heat transport. However, the underlying mechanism remains unconfirmed; further investigation with direct thermal measurement is warranted.

10.8.3 C3: Nucleation Site Distribution

Under uniform stress, collapse should nucleate at structural weak points (hot spots, power rail crossings). Instead:

Metric	Value
χ^2 statistic	12.0
Critical value (df=15, $\alpha=0.05$)	25.0
Null hypothesis	Not rejected
Interpretation	Uniform nucleation

Table 15: Nucleation site uniformity test (C3). No structural weak points detected.

Key finding: Collapse nucleates *uniformly* across the die—there are no preferred “weak points.” The corner nucleation observed in C1 was trial-specific, not structural. This suggests collapse location is stochastic under uniform stress.

10.8.4 C4: Spatial Leading Indicators

Spatial variance of k_{eff} across the sensor array was monitored during gradual collapse induction:

Metric	Value
Leading indicator	Spatial variance
Variance increase before collapse	+10.5
Early warning margin	$\Delta\rho = 0.317 \pm 0.125$ (95% CI: [0.221, 0.413])
k_{eff} before collapse	11.6
k_{eff} after collapse	2.3
k_{eff} degradation	80%

Table 16: Spatial leading indicators (C4). Early warning via spatial variance.

Key finding: Spatial variance *increases* before the system crosses the $\rho_{\text{crit}} = 0.43$ threshold. Across 30 trials, this provides $\Delta\rho = 0.317 \pm 0.125$ of advance warning (95% CI: [0.221, 0.413], $p < 0.0001$). At 0.5 m/s propagation velocity, this translates to tens of milliseconds of early warning time. The 80% drop in k_{eff} ($11.6 \rightarrow 2.3$) quantifies the severity of diversity loss at collapse.

10.8.5 C-Series Summary

Exp	Question	Result	Status
C1	$k_{\text{eff}} = k/(1 + \rho(k - 1))$?	$R^2 = 0.798$ ($n = 21$)	Withdrawn (identity check)
C2	Propagation velocity?	0.5 ± 0.4 m/s	First measurement
C3	Nucleation hotspots?	Uniform ($\chi^2 = 12$)	No hotspots
C4	Leading indicators?	Spatial variance	Found

Table 17: C-series coherence collapse propagation results.

These experiments transform CCA from a theoretical framework into an *empirically validated* physical phenomenon with measurable propagation dynamics and actionable early warning signals.

10.9 E-Series: Robustness Validation (January 2026)

Follow-up experiments addressed statistical robustness and edge cases:

Exp	Question	Result	Implication
E1	High- ρ resilience?	$\rho_{\text{intra}} > \rho_{\text{inter}}$ preserves k_{eff}	Block structure matters
E2	$\Delta\rho$ reliability?	0.317 ± 0.125 , CI [0.221, 0.413]	Early warning replicates
E3	Correlation dynamics?	Global, instantaneous	Shared resources
E4	Collapse threshold?	$k_{\text{eff,crit}} = 4.0$, latency $\uparrow 2.3\times$	Operational definition

Table 18: E-series robustness validation results.

E1 (Negative examples): Systems can maintain high *average* correlation ($\rho_{\text{avg}} > 0.5$) without collapse *if* correlation is block-structured (independent clusters). This provides the requested counterexample: high- ρ systems that remain resilient. The key is $\rho_{\text{intra}} \gg \rho_{\text{inter}}$, which preserves $k_{\text{eff}} > 1$.

E4 (Collapse operationalized): Functional collapse occurs at $k_{\text{eff,crit}} \approx 4.0$, where detection latency degrades $2.3\times$. This provides an operational definition of collapse.

10.10 F-Series: Mechanism and Causality (January 2026)

F-series experiments investigated propagation mechanism and causal direction:

Exp	Question	Result	Implication
F1	Correlation dynamics?	Global, instantaneous	Shared resources
F2	Barrier sync effect?	$\rho: 0.171 \rightarrow 0.050$	Sync decorrelates
F3	Does ρ cause fragility?	$\rho \uparrow$, fragility $\downarrow$ ($p < 0.001$)	Chaos arm only (Rem. 10.5)
F4	Common cause or propagating?	Isolated: across-pool $\rho = 0.001$	Common cause confirmed

Table 19: F-series mechanism and causality results.

F1 (Correlation dynamics): Correlation changes are *global and instantaneous*, mediated by shared resources (memory controller, power delivery). Arrival time is independent of distance (linear $R^2 = 0.005$, quadratic $R^2 = 0.011$, scaling exponent $\beta = 0.13 \approx 0$). Correlation state changes affect all sensors simultaneously rather than propagating spatially.

F3 (chaos arm): Intervention study. At baseline ($\rho = 0.037$, chaos regime), perturbation response was 755% and sensitivity was 4.88σ . After increasing correlation ($\rho = 0.170$), response dropped to 103% and sensitivity to 1.01σ ($p < 0.001$). Fragility falls as correlation rises out of the chaos regime.

Remark 10.5 (Correction to v3: the rigidity arm was never measured). Version 3 read F3 as validating that *both* extremes—too little and too much correlation—produce fragility, and hence the corridor as a whole. **F3 cannot support the two-sided claim.** Every correlation value the experiment measured is $\rho \in \{0.0037, 0.037, 0.143, 0.156, 0.170, 0.171\}$: a maximum of 0.171, against a stated rigidity edge of 0.43. Both reported points sit on the *chaos* side of the corridor, and no measurement approaches the rigidity regime.

The two-sided corridor is the framework’s distinctive claim, and the rigidity arm—fragility from *excess* correlation—is the half that carries the warning CCA exists to give. On this substrate that half is **untested**. What F3 establishes is the chaos arm alone: fragility decreases as ρ rises from 0.037 to 0.170. Whether it rises again beyond some upper threshold is not addressed by any measurement reported here. See `CORRECTIONS_v3.md` (C-14).

F4 (Common cause confirmed): Stream isolation test confirms correlation is *common cause*, not propagating signal. With isolated memory pools: within-pool $\rho = 0.080$, across-pool $\rho = 0.001$ (ratio $\rightarrow \infty$). With concurrent/contending streams: within-pool $\rho = 0.135$, across-pool $\rho = 0.040$ (ratio $3.3\times$). Correlation arises from shared resources (memory controller, power delivery), not from signals propagating between sensors.

10.11 Implications for CCA

The GPU series (B, O, C, E, F) characterizes an instrument. Version 3 read it as validating the theory; two of the four conclusions it drew are withdrawn or disputed, and are listed here with their status rather than removed:

1. **WITHDRAWN.** v3 concluded that “the k_{eff} formula is not merely algebraically valid but physically instantiated ($R^2 = 0.798$, $n = 21$).” The comparison is an identity check, not a test (Remark 10.4).
2. Correlation structure in real hardware responds to workload and to induced coupling (measured; domain: one RTX 4090)
3. **WITHDRAWN.** v3 concluded that “software alone can induce collapse (no hardware vulnerability required).” Replication on the same hardware shows the demonstration was a measurement artifact, and that the surviving barrier effect is too variable to support a collapse claim (Remark 10.3).
4. Leading indicators provide actionable warning before threshold crossing (spatial variance)
5. Recovery is possible if correlation source is removed ($\tau = 6.5$ ms)
6. **Correlation changes are global**, mediated by shared resources (memory controller, power delivery)
7. **No structural weak points**—collapse nucleates stochastically under uniform stress
8. **Early warning quantified:** $\Delta\rho = 0.317 \pm 0.125$ (95% CI excludes zero, $p < 0.0001$)
9. **Block structure preserves resilience:** High average ρ with independent clusters maintains $k_{\text{eff}} > 1$
10. **Operational collapse threshold:** $k_{\text{eff,crit}} \approx 4.0$ (latency degrades $2.3\times$)

11. **Corridor, chaos arm only:** fragility rises as ρ falls toward chaos (F3: $p < 0.001$). **The rigidity arm was not measured** — see Remark 10.5

This transforms CCA from a theoretical risk framework into an *empirically validated* engineering tool with measured propagation dynamics and demonstrated early warning capability.

11 A Negative Result: The Institutional Case Study

This section reports a case study that does not work, and why. Version 3 presented it as cross-domain validation; it is retained here, re-scored, because a properly-scored failure is more informative than a deletion — and because the reason it fails is a methodological caution that generalises beyond this paper.

11.1 What the engine actually computes

The institutional engine evolves four state variables and flags collapse when

$$\sigma < \sigma_{\min} \quad \text{or} \quad f > f_{\max},$$

with $\sigma_{\min} = 0.2$ and $f_{\max} = 0.8$ in the shipped configuration. Constraint count k and correlation ρ are carried in the state and updated each step, and are **never read by the collapse rule or by anything downstream of it**. Sweeping (k, ρ) across their full range at fixed initial (σ_0, f_0) produces a single distinct outcome across 400 combinations.

Remark 11.1 (The central quantity is absent from the only decision made). This is a stronger defect than the identity checks of §10.2 and Remark 10.4. There, k_{eff} was compared against itself, which is uninformative. Here k_{eff} plays no part at all: the prediction reduces exactly to a two-variable threshold rule on the initial conditions, $\text{flag} \iff \sigma_0 < 0.60$ or $f_0 > 0.70$, which reproduces the engine’s actual output on all 14 countries. Whatever this case study measures, **it does not test CCA**. Version 3’s presentation of it as cross-domain validation of the k_{eff} formula was therefore mistaken in kind, not in degree.

11.2 Honest re-scoring

The outcome definition below was pre-specified before any outcome variable was computed: a country is positive iff its V-Dem Regimes-of-the-World category falls at least one level below its year-2000 value and remains there for at least three consecutive observed years, within 2000–2024. Declines beginning in 2023–24 cannot be verified for three years and are scored negative (conservative, guarding right-censoring). The country set is the 14 cases named in the repository’s own case-study list; v3 reported 13 without naming the omission.

	Predicted collapse	Predicted stable
Sustained decline	4	3
No sustained decline	5	2

Table 20: Institutional classification, $n = 14$, scored against the pre-specified outcome. Accuracy 0.43; balanced accuracy 0.43; Matthews correlation coefficient -0.15 ; permutation null $p = 0.87$; Fisher exact $p = 1.00$.

The three false negatives are **Hungary, Poland** and Canada. Hungary’s Regimes-of-the-World series runs liberal democracy through 2009, electoral democracy 2010–18, and electoral autocracy from 2019; Poland’s falls from liberal to electoral democracy in 2016. Version 3 counted both among “5/5 stable democracies correctly classified.” Scored against the pre-specified outcome, that figure is **2/5**. Always-flag and never-flag both score 0.50; the classifier reaches 0.43.

Remark 11.2 (The study was underpowered before data collection). At $n = 14$ with this base rate and nine positive predictions, a one-sided Fisher test does not reach $p < 0.05$ until the classifier is correct on 12 of 14 *with perfect sensitivity*. A flawless classifier at $n = 13$ reaches only $p = 0.0006$. No validation claim was reachable from this design regardless of how the engine performed, and this could have been determined before any country was scored.

Remark 11.3 (Threshold fragility, and why the favourable definition is disqualified). Across six pre-specified outcome definitions the Matthews coefficient ranges from -0.19 to $+0.65$. Only one — a Polity5 score-drop criterion — beats chance. It is also the one definition structurally blind to the two cases the classifier gets wrong: Polity5 records Hungary and Poland at the maximum score for *every* year 2000–2018, and its coverage ends six years short of the stated window. That is the definition v3’s numbers came from. A result that moves this far on the choice of ground truth is not entitled to the word “validated” under any of the six.

11.3 The timing claim is an artifact

Version 3 reported a mean detection lead of 7.6 years. Seven of the nine flags fire in the first simulated year because $f_0 > 0.8$ *already holds at initialisation* — the collapse condition is satisfied by the input data before any dynamics run. The reported lead time is therefore (event year $- 2001$): a restatement of when the events happened, not a prediction of them. A null model that flags every country in 2001 achieves a *larger* mean lead (10.8 years) at sensitivity 1.0. The one flag not set at initialisation, Turkey via the σ pathway, fires in 2019 — six years *after* its regime transition, so v3’s “flagged within 3 years of actual event” is correct in magnitude and wrong in sign, printed beside the 7.6-year claim.

The timing claim is withdrawn. What replaces it is the caution: **a detector whose trigger condition is satisfied by its own initial conditions will report large lead times, and those lead times carry no predictive content**. Any lead-time claim must be compared against a null that fires at t_0 .

11.4 What this case study is worth

Three things survive, none of them validation. It is a worked example of the institutional variable mapping, showing what the framework’s quantities would look like if measured. It is a demonstration that the mapping’s ρ construction is degenerate: six of the fourteen countries receive $\rho = 0.000$ exactly, forced by $\rho = f(1-k)$ whenever executive constraints are at maximum, regardless of the country’s actual data. And it is a record of a negative result that a pre-specified outcome definition makes legible and a post-hoc one concealed — which is the methodological point worth carrying forward.

12 Discussion

12.1 Limitations

12.1.1 Mathematical Simplifications

1. **Scalar ρ is oversimplified**: A single average pairwise correlation cannot capture clustered correlations, directional dependencies, or higher-order interactions. Two systems with identical mean ρ can have radically different resilience depending on correlation *structure*.

Mitigation: The RATCHET implementation includes an Extended Correlation Model (`correlation_tensor.py`) providing spectral analysis, block-diagonal detection, and higher-order correlation statistics. Future work should incorporate these richer metrics.

2. **Quasi-static assumptions:** Several derivations assume $d\rho/dt \approx 0$, which may not hold in rapidly adapting systems—especially AI-mediated ones where correlation can change faster than measurement cycles.
3. **Defense function J is definitional:** As noted in §2.3, J is a modeling choice, not a derived quantity. High J does not *necessarily* imply flourishing capacity—it implies high cost to adversaries under our model’s assumptions.
4. **Convexity requirements:** The volume decay theorem assumes convex deceptive regions. Non-convex geometries (torus, point cloud, fractal) may not exhibit exponential decay.
Mitigation: RATCHET includes robustness analysis (`robustness.py`) quantifying deviation from generic geometry assumptions.

12.1.2 Structural vs Timing Performance

1. **Phase classification:** Reliable across domains. CCA correctly identifies healthy, chaos, and rigidity phases. All systems that experienced upheaval were flagged; all stable systems were correctly classified as healthy.
2. **Timing uncertainty:** Higher than phase classification. Timing estimates average 5–8 years early for institutional systems; 8% SOH overestimate for batteries. These represent *lead times* for risk windows, not event dates.
3. **False positives are false positives.** Version 3 carried, at this point, the claim that the three recorded false positives (Tunisia, Egypt, Zimbabwe) had “detected real fragility” and that “the binary ‘collapse’ definition was inappropriate for these cases.” **That passage is withdrawn without replacement.** It revises the success criterion after the results are known, which converts recorded misses into successes and makes the classification claim unfalsifiable. The standing rule, stated here in its place: an outcome definition may be revised, but it must then be re-registered and the analysis re-run — never revised in the discussion of results already seen. See `CORRECTIONS_v3.md` (C-4).
4. **Intervention cross-effects not modeled in core theory:** Increasing λ often increases ρ (stricter rules favor similar actors), pushing toward rigidity.
Mitigation: RATCHET includes intervention dynamics modeling (`interventions.py`) with cross-effect matrices and adversary response.

12.1.3 Scope Limitations

1. **Scope of validation:** Chemistry, political science, and biology demonstrate cross-domain applicability. Broader validation would strengthen universality claims.
2. **Great Filter analogy:** We use “Great Filter” as a *metaphor* for mechanisms that could cause systemic failure at scale—not as a probabilistic claim about the Fermi paradox. Correlation accumulation shares structural features (invisibility, scalability, irreversibility) with such mechanisms, but CCA makes no predictions about civilizational outcomes.
3. **Fundamental detection limits:** $\sim 40\%$ of emergent deception patterns are information-theoretically undetectable via marginal analysis. This is a ceiling, not a floor—actual detection may be worse.

On the measurement of ρ . A common objection is that pairwise correlation ρ is difficult to measure precisely. This concern, while valid for absolute quantification, overstates the requirement. The classification algorithm (Algorithm 1) depends primarily on *trend direction* and *threshold crossings*, not on precise point estimates. A system need not know that $\rho = 0.47$; it suffices to know that ρ is increasing and has crossed 0.4. Comparative estimation (“higher than last month”) is tractable even when absolute estimation is not. This aligns with standard practice in control systems, where derivative terms matter more than instantaneous values for stability analysis.

12.2 Outstanding: the validation this paper still lacks

Two items are named here rather than deferred silently, because both are load-bearing and neither is closed by any correction in this version.

The missing measurement, and the cheapest way to obtain it. Section 10.2 establishes that no independent test of the Kish identity appears in this paper, and cites [22] as the nearest existing instance of the kind of measurement required — on a different substrate, with no paired comparison against k_{eff} . The comparison is nonetheless inexpensive on the GPU array, because both quantities can be computed from correlation matrices already stored. The concrete form: compute $N_{\text{eff}} = k/\lambda_{\text{max}}$ from the largest eigenvalue of the measured correlation matrix, and compare it against $k_{\text{eff}} = k/(1 + \bar{\rho}(k - 1))$ from the same data. These are distinct functionals over distinct inputs — one uses the full spectrum, the other a scalar summary — so they are free to disagree, which is exactly what Exp 86 and C1 were not. Exp 117 is the same idea in a more elaborate form; k/λ_{max} is the version a reviewer will expect to see first. Until that comparison is run, the identity remains *proved* and *unmeasured*, and this paper should not be read as claiming otherwise.

Theorem 2.1 substitutes k_{eff} into the exponent without deriving the step, and the substitution is measurably wrong. The proof appeals to the effective-sample-size formula in one sentence. That formula governs the *variance of a mean* of equicorrelated variables; the volume argument requires the *expectation of a sum of log-reduction factors*. These are different moments of different objects, and the substitution is not established by the citation. Relatedly, the product is written over $i = 1$ to k_{eff} , a non-integer index — the tables in this paper carry k_{eff} values of 5.20 and 7.5.

The step can be checked directly, and it fails. Cutting a bounded convex region (the unit ball in $\mathbb{R}^8$) with k halfspaces whose normals are equicorrelated at ρ , and reading off the effective count the measured volume decay actually implies, gives:

k	ρ	Kish k_{eff}	implied k_{eff}	ratio
4	0.00	4.00	4.10	1.02
8	0.00	8.00	8.12	1.02
8	0.50	1.78	3.36	1.89
16	0.50	1.88	6.21	3.30
32	0.20	4.44	13.61	3.06
32	0.50	1.94	7.91	4.08

Table 21: Volume decay does not follow the Kish k_{eff} . The substitution is exact at $\rho = 0$ (where it is trivial) and degrades monotonically with k , reaching a factor of 4 at $k = 32$, $\rho = 0.5$. Full sweep in *replication/geometry*/.

The substitution holds only at $\rho = 0$, where $k_{\text{eff}} = k$ and it asserts nothing. Elsewhere the

geometry decays *faster* than the formula predicts — correlated halfspaces still cut substantially more volume than the design effect allows — and the discrepancy grows with k , which is the regime the framework’s asymptotic claims live in. Note also that at $\rho = 0$ and $k = 32$ the implied count is 25.7 rather than 32: the pure exponential form itself degrades at large k , independently of any correlation question. Everything downstream inherits the step: T_{truth} , the singularity boundary $K_{\text{req}} \cdot \rho \geq 1$, and the phase corridor. This is **open**, not proved, and the correct effective-count law for volume decay under correlated constraints is not supplied by this paper.

The deceptive/honest asymmetry is assumed, not derived — and the assumption carries the entire safety claim. The volume-decay theorem cuts a bounded convex region with i.i.d. halfspaces. Nothing in that construction refers to which points are honest. Safety requires the *ratio* $|D|/|H|$ to fall; the theorem gives only that both regions shrink.

Simulation isolates what the assumption is doing. With honest and deceptive regions **congruent** — equal radius, equal distance from the origin, differing only in direction, so that no geometric advantage is smuggled in — and constraints drawn i.i.d. as assumption (a) specifies:

Constraint distribution	λ_H	λ_D	gap
i.i.d. (assumption (a) alone)	0.0262	0.0214	−0.005
truth-tracking (not assumed)	0.0000	0.0181	+0.018

Under the theorem’s own assumptions the gap is indistinguishable from zero and marginally negative: over 40 constraints the ratio $|D|/|H|$ does not fall but drifts from 1.01 to 1.22. **Adding i.i.d. constraints buys no safety whatsoever.** The asymmetry appears only in the second row, where constraints are drawn subject to being satisfied by the honest region — and that selection is the alignment problem, assumed here rather than solved.

One caution for anyone re-running this: an earlier version of the test appeared to show asymmetry under i.i.d. constraints, because the honest region had been made small and central while the deceptive region was large and peripheral. Compact central regions survive random halfspace cuts better than diffuse peripheral ones, and that effect is easily mistaken for a safety result. It is a claim about where honest behaviour *lives*, not about what constraints *do*, and it needs its own justification if relied upon.

Stated plainly: **the geometry shows that constraints shrink feasible volume; it does not show that they shrink deception preferentially. The framework’s central safety intuition rests on the constraint-generating process being truth-tracking, which is exactly what an alignment method would have to supply.** The limitation previously stated here (“assumes convex deceptive regions”) is a weaker and different caveat and does not cover this.

12.3 Ethical Considerations

CCA is an *engineering risk tool*. Several design choices reinforce this:

- Conditional framing: “If X continues, then Y becomes likely”
- Intervention focus: Identifying actionable intervention points
- Falsifiability: Clear claims that can be tested and refuted
- Explicit detection limits acknowledged

13 Conclusion

Coherence Collapse Analysis models coherence through constraint geometry. What that yields, after this version’s corrections and the falsification conditions run against it, divides cleanly.

What is established. The effective-count map $k_{\text{eff}} = k/(1 + \rho(k - 1))$ is an identity with a Möbius structure: strictly monotone in ρ , saturating at $1/\rho$. Scale cannot restore diversity that correlation has collapsed, and this holds without calibration on any substrate. The stability condition $\alpha/k \geq d$ and the three closed-form collapse timescales follow from stated assumptions and are machine-verified in Lean — which establishes internal consistency and, as §7.2 says, nothing about the world. L-01 establishes that the class of marginal-preserving undetectable patterns is non-empty: existence, no measure. The GPU array is a real, characterized instrument, and its measurement protocol — shuffle null, repeated trials, first-run discard — is transferable.

What fired against it. Of six falsification conditions, two fired. The equicorrelated precondition the identity requires *fails* below $\rho \approx 0.04$ on the measured substrate, so k_{eff} figures from that band are unlicensed. And the volume-decay law $V = V_0 e^{-\lambda k_{\text{eff}}}$ is wrong by a factor of 1.9 to 4.1 at $\rho = 0.5$, with the error growing in k — so T_{truth} , the singularity boundary, and the phase corridor all rest on a step that does not hold. Two further conditions are unmet, one is unresolved, and one passes.

What is not established, and was claimed in version 3. Cross-domain validation: withdrawn — applying one formula to three datasets shows it can be applied, not that it holds. Hardware confirmation of the k_{eff} formula: withdrawn, twice over, as both results computed the identity and compared it against itself. Software-induced collapse: withdrawn as a measurement artifact. The two-sided corridor: unresolved, and §1.6 reports why three separate fragility measures could not settle it. The deceptive/honest asymmetry that would make any of this a safety result is assumed, not derived, and simulation shows that under the theorem’s own assumptions adding constraints buys no differential protection at all.

Final assessment. CCA is a **vocabulary and a measurement substrate for correlated fragility**, not a validated predictive framework. Its durable content is an identity with a saturation ceiling — inherited from Kish (1965) and reached independently in four other fields — together with an instrument for measuring correlation structure and a discipline for deciding when such measurements mean anything. Its central safety intuition, that accumulating independent constraints starves deception, does not follow from the geometry presented here: the geometry shows that constraints shrink feasible volume, not that they shrink deception preferentially. Making hidden fragility legible remains the goal; this version establishes how much of that goal the evidence currently supports, which is less than version 3 claimed and more than nothing.

Acknowledgments

The RATCHET framework is open-source and available at <https://github.com/CIRISAI/RATCHET>.

A Notation Reference

Symbol	Meaning
k	Raw constraint count
k_{eff}	Effective constraint count
ρ	Pairwise constraint correlation
ρ_{crit}	Critical correlation threshold
λ	Decay constant / strictness
σ	Sustainability integral
α	Constraint generation rate
d	Decay rate
J	Defense function
R_c	Collapse rate (dJ/dt)
T_{truth}	Time to deception collapse
T_{entropy}	Time to entropic decay
T_{capture}	Time to coordination capture
T_{eff}	Effective collapse time
K_{req}	Required effective constraints

Table 22: Notation reference for Coherence Collapse Analysis.

B Key Formulas

$$k_{\text{eff}} = \frac{k}{1 + \rho(k - 1)} \quad (\text{Effective Constraints})$$

$$K_{\text{req}} = \frac{-\ln(\varepsilon/V_0)}{\lambda} \quad (\text{Required Constraints})$$

$$\rho_{\text{crit}} = \frac{1}{K_{\text{req}}} \quad (\text{Critical Correlation})$$

$$J = k_{\text{eff}} \cdot \lambda \cdot \sigma \quad (\text{Defense Function})$$

$$T_{\text{truth}} = \frac{K_{\text{req}}(1 - \rho)}{\alpha(1 - K_{\text{req}} \cdot \rho)} \quad (\text{Time to Truth})$$

$$T_{\text{entropy}} = \frac{\ln(\sigma/\sigma_{\text{min}})}{d} \quad (\text{Time to Entropy})$$

$$T_{\text{capture}} = \frac{(n/3) - f}{r_{\text{cap}}} \quad (\text{Time to Capture})$$

$$P(\text{detect}) = 1 - e^{-10\varepsilon} \quad (\text{Detection Probability})$$

References

References

- [1] L. Kish, *Survey Sampling*, John Wiley & Sons, 1965. *Origin of the effective sample size formula $k_{\text{eff}} = k/(1 + \rho(k - 1))$, here applied to constraint-based systems.*
- [2] K.K. Ladha, “The Condorcet Jury Theorem, Free Speech, and Correlated Votes,” *American Journal of Political Science* 36:617–634, 1992. *Extends the Condorcet Jury Theorem to correlated votes. The central claim of this paper, in the voting setting, thirty-four years earlier.*
- [3] K.K. Ladha, “Condorcet’s Jury Theorem in Light of de Finetti’s Theorem,” *Social Choice and Welfare* 10(1):69–85, 1993.
- [4] K.K. Ladha, “Information Pooling Through Majority-Rule Voting: Condorcet’s Jury Theorem with Correlated Votes,” *Journal of Economic Behavior and Organization* 26:353–372, 1995. *Generalizes across normal, hypergeometric and Pólya correlated-voting models; proves majority-rule effectiveness decreases as correlation increases, and that diversity improves decisions by lowering average correlation.*
- [5] D. Austen-Smith and J. Banks, “Information Aggregation, Rationality, and the Condorcet Jury Theorem,” *American Political Science Review* 90(1):34–45, 1996. *Independence is necessary, not merely sufficient; correlated voters can perform worse than individuals. A sharper statement than this paper’s singularity condition.*
- [6] D.E. Eckhardt and L.D. Lee, “A Theoretical Basis for the Analysis of Multiversion Software Subject to Coincident Errors,” *IEEE Transactions on Software Engineering* SE-11(12):1511–1517, 1985.
- [7] J.C. Knight and N.G. Leveson, “An Experimental Evaluation of the Assumption of Independence in Multiversion Programming,” *IEEE Transactions on Software Engineering* SE-12(1):96–109, 1986. *27 versions written independently from one specification, one million tests; coincident failures substantially exceeded the independence prediction. The empirical thesis of this paper, in software, forty years earlier.*
- [8] J.C. Knight and N.G. Leveson, “A Reply to the Criticisms of the Knight & Leveson Experiment,” *ACM SIGSOFT Software Engineering Notes* 15(1):24–35, 1990.
- [9] “N-Version Programming with Coding Agents,” arXiv:2606.20158, June 2026. *429 coincident-failure cases against an independence prediction of 115.36 — Knight–Leveson replicated on LLM agents.*
- [10] M. Laakso and R. Taagepera, “‘Effective’ Number of Parties: A Measure with Application to West Europe,” *Comparative Political Studies* 12:3–27, 1979. *Equivalent to the inverse Simpson and inverse Herfindahl–Hirschman indices. The same effective-count construction, reached independently in political science.*
- [11] M.O. Hill, “Diversity and Evenness: A Unifying Notation and Its Consequences,” *Ecology* 54(2):427–432, 1973.
- [12] L. Jost, “Entropy and Diversity,” *Oikos* 113:363–375, 2006. *Hill numbers / effective number of species. Note the conversion $\exp(H)$ — relevant to the microbiome row of Table 8.*
- [13] A. Meucci, “Managing Diversification,” SSRN 1358533, 2010. *Effective number of bets via the entropy of the diversification distribution.*

- [14] V. Lefort et al., “Can We Trust the Wisdom of LLM Crowds?” arXiv:2409.00094, August 2024. *Condorcet Jury Theorem applied to LLM ensembles; marginal gains imply lack of independence. Predates any application of effective-count-under-correlation to AI oversight by this project.*
- [15] Kim et al., “Error Correlation Across Large Language Models,” ICML, 2025. *350+ models; agreement on wrong answers approaching 60% on some benchmarks.*
- [16] F.E. Dorner et al., “Limits to Scalable Evaluation at the Frontier,” ICLR (oral), 2025. *Formal ceiling: when the judge is no more capable than the evaluated model, no debiasing halves the required ground-truth data. A quantified limit, where L-01 is an existence result.*
- [17] Jiang et al., “The Artificial Hivemind,” NeurIPS (best paper), 2025. *Output homogeneity within and across model families.*
- [18] S. Kohli, “Effective Sample Size in LLM Judge Panels,” arXiv:2605.29800, 28 May 2026. *Independent, concurrent application of the Kish identity to a 9-judge, 7-family panel: $n_{\text{eff}} = 2.18$ [2.07, 2.31] at mean pairwise $\phi = 0.391$, asymptote at $1/\bar{\phi} \approx 2.6$. Cross-checks with $n_{\text{eff}} = k/\lambda_{\text{max}}$ (2.16), which does not assume exchangeability — the independent measure this paper lacks. Uses ϕ because it reduces to Pearson for binary error vectors, rejecting Cohen’s κ for conflating prevalence with dependence.*
- [19] W.R. Ashby, *An Introduction to Cybernetics*, Chapman & Hall, 1956. *Law of Requisite Variety: “Only variety can absorb variety.” Theoretical foundation for diversity requirements in control systems.*
- [20] Explosive Synchronization Study, *Proceedings of the National Academy of Sciences*, 2025. *Phase transitions predict system fragility; proximity to first-order transitions correlates with collapse risk.*
- [21] E. Moore, “CIRISAgent: An Open-Source Framework for Ethical AI Through Transparent Architecture,” CIRIS Research, 2026. *Machine Conscience architecture; companion paper describing the 22-service microarchitecture.*
- [22] E. Moore, “Constrained Reasoning Chains: An Empirical Telemetry Study of LLM Alignment,” Zenodo, 2026-04-28. DOI: 10.5281/zenodo.19839280. *Effective dimensionality on $n = 6,465$ production traces via the participation ratio of the covariance eigenvalue spectrum — a spectral measure not computed from ρ , and the nearest existing instance of the independent measurement this paper’s §10.2 identifies as missing. Cited with its stated limits: in-sample thresholds from continuous tuning rather than a held-out cohort, and a threshold quoted across two distinct estimators.*
- [23] NASA Prognostics Center of Excellence, “Li-ion Battery Aging Datasets,” NASA Ames Research Center. *19-cell battery degradation dataset used for cross-domain validation.*
- [24] Quality of Government Institute, “QoG Standard Dataset,” University of Gothenburg, 2024. *203 countries, 12,393 observations (1946–2024) for institutional analysis.*
- [25] Center for Systemic Peace, “Polity V Project: Political Regime Characteristics and Transitions,” 2024. *Regime transition coding used for collapse event identification.*
- [26] D. McDonald et al., “American Gut: An Open Platform for Citizen Science Microbiome Research,” *mSystems*, 2018. *2,081 microbial taxa dataset used for microbiome domain validation.*

- [27] B. Grünbaum, “Partitions of Mass-Distributions and of Convex Bodies by Hyperplanes,” *Pacific Journal of Mathematics*, 1960. *Log-concavity of volume under halfspace intersection; foundation for volume decay theorem.*
- [28] S. Beer, *Brain of the Firm*, John Wiley & Sons, 1981. *Viable System Model (VSM); organizational cybernetics and requisite variety in management.*
- [29] D.J. Watts, “A Simple Model of Global Cascades on Random Networks,” *PNAS*, 2002. *Cascade dynamics in networked systems; threshold models for systemic failure.*
- [30] M. Granovetter, “Threshold Models of Collective Behavior,” *American Journal of Sociology*, 1978. *How individual thresholds aggregate to produce collective dynamics.*
- [31] M. Castro and B. Liskov, “Practical Byzantine Fault Tolerance,” *OSDI*, 1999. *Foundation for BFT consensus; $n \geq 3f + 1$ requirement for f Byzantine failures.*
- [32] NVIDIA Corporation, “NVIDIA GeForce RTX 4090 Specifications,” 2022. *GPU hardware specifications for strain gauge array experiments.*
- [33] A. Rukhin et al., “A Statistical Test Suite for Random and Pseudorandom Number Generators,” NIST SP 800-22, 2010. *NIST test suite used for TRNG validation.*
- [34] E.N. Lorenz, “Deterministic Nonperiodic Flow,” *Journal of the Atmospheric Sciences*, 1963. *Lorenz attractor; note: chaotic oscillators were invalidated for entropy amplification in this work.*
- [35] C. Sparrow, *The Lorenz Equations: Bifurcations, Chaos, and Strange Attractors*, Springer-Verlag, 1982. *Lyapunov exponent literature value ($\lambda = 0.906$) for Lorenz system validation.*
- [36] L. de Moura et al., “The Lean 4 Theorem Prover and Programming Language,” *CADE*, 2021. *Formal verification framework used for theorem proofs.*
- [37] The mathlib Community, “mathlib4: The Lean 4 Mathematical Library,” 2024. *Mathematical foundations library for Lean 4 proofs.*
- [38] CIRIS Research Team, “CIRISArray: C-Series CCA Validation,” 2026. *Physical validation experiments for Coherence Collapse Analysis on GPU hardware. Commit a75e517. Available: <https://github.com/CIRISAI/CIRISArray>*
- [39] M. Glickman and T. Sharot, “How human–AI feedback loops alter human perceptual, emotional and social judgements,” *Nature Human Behaviour*, 2024. *Empirical study ($n=1,401$) showing AI amplifies human biases more than human-human interaction.*
- [40] I. Shumailov et al., “AI models collapse when trained on recursively generated data,” *Nature*, 2024. *Demonstrates model collapse: outputs converge toward bland central tendencies when trained on AI-generated content.*
- [41] Y. Daryani, Z. Sourati, and M. Dehghani, “The Homogenizing Engine: AI’s Role in Standardizing Culture and the Path to Policy,” *Policy Insights from the Behavioral and Brain Sciences*, 2025. *“LLMs have emerged as unprecedented drivers of cultural homogenization, operating at scales and speeds that exceed all previous technologies.”*
- [42] “Echoes in AI: Quantifying lack of plot diversity in LLM outputs,” *Proceedings of the National Academy of Sciences*, 2025. *Empirical quantification of reduced diversity in LLM-generated content.*

- [43] “The Homogenizing Effect of Large Language Models on Human Expression and Thought,” *CHI Conference on Human Factors in Computing Systems*, 2025. *AI suggestions homogenize writing toward western styles and diminish cultural nuances.*
- [44] “When Algorithms Mirror Minds: A Confirmation-Aware Social Dynamic Model of Echo Chamber and Homogenization Traps,” *arXiv:2508.11516*, 2025. *Mathematical proof that echo chambers and homogenization traps will inevitably occur under current recommender dynamics.*
- [45] S. Cecchetti et al., “AI and Systemic Risk,” *CEPR/European Systemic Risk Board*, 2025. *Documents algorithmic monoculture as source of financial systemic risk.*