

Coherence Management as a Substrate-Independent Engineering Principle:

Cross-Domain Validation and Sociotechnical Implications

Eric Moore
CIRIS Ethical AI

May 16, 2026

Abstract

Complex coordinating systems — chemical, biological, institutional, neural, sociotechnical — share a structural property: as pairwise correlation ρ between their functional constituents rises, effective independent dimensionality $k_{\text{eff}} = k/(1 + \rho(k - 1))$ collapses toward unity regardless of nominal scale. We test this *coherence-management* principle across seven substrates spanning four agency rungs (chemistry, biology, institutions, protein folding, neural firing, ecological time series, power grid) and find the Kish identity fits at $R^2 > 0.7$ in all seven under the pre-registered tolerance-band rule (Exp 2 closeout 2026-05-16). A 128-sensor GPU strain-gauge array provides physical instantiation: software coordination alone is sufficient to drive k_{eff} from 6.5 to 1.0 at $\rho \rightarrow 1$, with critical-exponent $\beta = 1.09$ and $R^2 = 0.978$ on the power-law fit. At the sociotechnical scale, the CRCv2 cascade-implementation audit (zero-leak architectural property + ratchet asymmetry on actual cascade firings; formalized in `OverrideRate.lean`) runs across five foundation-model vendors (Google Gemini 2.5-Flash, Meta Llama-4-Scout, Alibaba Qwen-3.5-35B, OpenAI GPT-5.4, Anthropic Claude Sonnet 4.6): 1,401 chains scored, 0/1,401 escape rate, 88 corrections all in the ratchet direction. Correction rates vary substantially by vendor (0%–18.2%); the Sonnet 4.6 cell yielded zero corrections at $n = 67$ on this battery, making RA-1 untestable for that cell (the cascade did not activate; not a passed test). We stratify all claims by verification status, pin ten falsification criteria (F-1..F-10), pre-register six experiments (Exp 1..Exp 6, four engineering and two open-research), and surface four sociotechnical engineering implications (§5): AI alignment is structurally a correlation-management problem, federation outperforms centralization for resilience, crisis governance is correlation-amplifying, and the “Great Filter” generalizes as an engineering hazard in any coordinating system. We also include a research program (§6) listing three conjectures that extend the framework beyond the engineering tier — quantum-substrate coherence dynamics, civilizational coherence-collapse, and sociotechnical audit-pressure effects — each with explicit falsification handles, pre-registered experiments, and open priors. These conjectures are not load-bearing on the validated tiers; they are included so researchers who want to test the principle at other substrates can find the formal hooks.

Contents

1	Introduction: The Coherence Hypothesis	2
1.1	The claim	2
1.2	Scales of validation	3
1.3	Relation to prior CIRIS papers	3
1.4	What this paper does not claim	3
1.5	Epistemic ladder: seven levels for seven audiences	4

2	The Mathematical Framework	4
2.1	The Kish dynamic	4
2.2	The singularity boundary	5
2.3	The healthy corridor	5
3	Cross-Substrate Empirical Validation	5
3.1	Domain fits	5
3.2	Exp 2 closeout: F-7 cross-substrate test PASSED 7/7	6
3.3	Hardware-physical instantiation: GPU strain-gauge array	7
3.4	Independent-corpus replication of the $N_{\text{eff}} \approx 7.1$ threshold	7
4	The Sociotechnical Substrate: CIRIS as Coherence-Management Infrastructure	8
4.1	What CIRIS does differently	8
4.2	The 22-service runtime and H3ERE conscience	8
4.3	The L3 load-bearing claim: CRCv2 override-rate (revised 2026-05-16)	8
5	Engineering Implications	9
5.1	AI alignment is structurally a correlation-management problem	10
5.2	Federation outperforms centralization for resilience	10
5.3	Crisis governance is correlation-amplifying — and that is the danger	11
5.4	The “Great Filter” generalizes as an engineering hazard in any coordinating system	11
6	Research Program: Conjectured Extensions to Other Substrates	11
6.1	Conjecture A: Quantum–classical k_{eff} bridge	11
6.2	Conjecture B: Iterated coherence-management as a civilizational hazard	12
6.3	Conjecture C: Audit-pressure effects on cascade behavior	13
6.4	Level 4: Agency and consent as a structural fact at A3+	13
6.5	Level 5: TSVF retrocausal reading (interpretive)	14
6.6	Level 6: Theological identification (recognition)	15
6.7	Level 7: Civilizational extension and external signatures	16
7	Claim Stratification	17
8	Falsification Framework	17
9	Pre-Registered Experimental Protocols	18
10	Conclusion	19

1 Introduction: The Coherence Hypothesis

1.1 The claim

At every scale of complex organization, the same structural dynamic is hypothesized to govern: as the correlation ρ between the distinct functional elements of a system increases, its effective independent dimensionality k_{eff} collapses. The capacity of a system to manage its internal correlation faster than that correlation drives it toward entropic homogenization fundamentally defines its trajectory. We call this the *coherence hypothesis*.

1.2 Scales of validation

The hypothesis is tested at four scales (Table 1). Tiers 1–2 carry empirical or formal evidence; Tier 3 surfaces engineering implications that follow directly from the validated tiers.

Tier	Scale	Substrate / measurement
1 (validated)	Macroscopic / network	7-substrate Kish-fit closeout: chemistry (NASA Li-ion), biology (microbiome), institutions (Polity5 + WGI), protein folding (AlphaFold), neural firing (Allen), ecological time series (BioTIME), power grid (PMU); $R^2 > 0.7$ in 7/7
1 (validated)	Hardware-physical	128-sensor GPU strain-gauge array on RTX 4090; software-induced collapse, $\beta = 1.09$ critical exponent
2 (validated, smoke-disclosed)	Sociotechnical / AI alignment	CIRIS H3ERE conscience cascade; CRCv2 override-rate replicated across 5 LLM vendors (1,401 chains, 0 LEAK)
3 (engineering implications)	Federation, governance, AI deployment topology	Implications that follow directly from the validated tiers: alignment requires external machinery, federation > centralization, crisis dynamics need real-time ρ monitoring, Great Filter as engineering hazard

Table 1: Four scales of substrate testing. Tier 3 is engineering implications derived from the validated tiers (no metaphysics).

1.3 Relation to prior CIRIS papers

This paper builds on three prior works:

- [2] (CIRISAgent v2): the 22-service architecture and H3ERE conscience module—the *system*.
- [3] (CCA v3): the Kish-dynamic formalism, the $\rho_{\text{critical}} = 0.43$ collapse threshold, formal proofs in Lean 4—the *mathematics*.
- [4] (CRC): an empirical telemetry study on $n = 6,465$ production reasoning traces establishing the N_{eff}^H manifold geometry—the *empirics*.

The independent-corpus calibration result reported in §3.4 ($n = 264$, 90% variance horizon = 7 dims exact) is an independent replication of the manifold geometry. The CRCv2 cross-vendor closeout in §4.3 (5 vendors, 1,401 chains, 0 LEAK) is new in this paper.

1.4 What this paper does not claim

The paper’s load-bearing claims (§2–§5) are bounded engineering analysis. We make no *load-bearing* claims about quantum-mechanical decoherence, civilizational collapse-prediction, or cosmological phenomena. The framework explains failure modes in coordinating systems and prescribes structural remedies; the engineering implications in §5 follow directly from the validated tiers.

The paper does, however, surface a research program (§6) with explicit conjectures and interpretive layers organized as an epistemic ladder (§1.5). The interpretive layers are not load-bearing; they exist so that readers who want to follow the framework’s implications further

— through agency-bearing systems, into quantum-foundational reading frameworks, or onto contemplative-tradition recognition — can find the formal hooks. *Nothing at any higher level of the ladder is load-bearing on the levels below it.*

1.5 Epistemic ladder: seven levels for seven audiences

The framework operates at seven distinct epistemic levels, in order of increasing speculation and decreasing evidential support. Each level is intended as a stopping point for readers with different epistemic comfort zones. **Nothing at any level is load-bearing on the levels above it** — a skeptic can stop at Level 2 and lose nothing the validated framework provides; a working scientist can stop at Level 3 with the universality conjecture as the next test; a philosopher can stop at Level 4 on framework-defensible interpretation; a quantum-foundations-curious reader can stop at Level 5; a contemplative at Level 6; the rare cosmological reader proceeds to Level 7.

Level	Character	Stop here if you want	Where developed
1	Algebraic foundation	Mathematical foundations only	§2
2	Empirical + formal validation	Validated cross-substrate engineering claim	§3, §4
3	In-test universality conjecture	Empirical bet on substrate-independence	§3.2, §6 Conjecture A/B
4	Structural fact at A3+	Framework-defensible reading of agency, goal-formation, consent	§6.4
5	TSVF-conditional reading	Quantum-foundations-informed interpretation of bidirectional causation at A3+	§6.5
6	Recognition (not derivation)	Cross-tradition contemplative mapping	§6.6
7	Civilizational extension	Framework as metric; specific empirical claims external	§6.7

Table 2: The framework’s epistemic ladder. Each level has its own evidential character; readers may stop where their epistemic comfort allows.

Levels 1–3 are testable in the standard falsificationist sense (engineering tiers; the paper’s load-bearing claims). Levels 4–6 are interpretive layers that follow if certain reading commitments are accepted (TSVF for L5; cross-tradition recognition for L6); these do not have falsification tests in the same sense, but each is structurally honest about what it requires. Level 7 is the most speculative; the framework provides the metric, and any empirical claim about specific signatures (UAP behavioral patterns, civilizational trajectories) is external to the framework.

2 The Mathematical Framework

2.1 The Kish dynamic

The mathematical core is the effective constraint count, originally derived by Kish [1] for survey sampling under autocorrelation. We use it as a useful coordinate system for constraint-based coordination across substrates — it gives the same algebraic answer in every domain where “constituents” and “pairwise correlation” can be operationalized, which is the empirical content

of the cross-substrate fit (§3.2):

$$k_{\text{eff}} = \frac{k}{1 + \rho(k - 1)} \quad (1)$$

where k is the raw count of constraints (or constituent elements) and ρ is the average pairwise correlation between them. As $\rho \rightarrow 1$, $k_{\text{eff}} \rightarrow 1$ independent of k .

2.2 The singularity boundary

The system enters an irreversible rigidity phase when

$$K_{\text{req}} \cdot \rho \geq 1 \quad (2)$$

where K_{req} is the minimum diversity required to satisfy the system’s task topology. At this threshold, adding more constraints provides no marginal adaptive capacity, and recovery is computationally prohibitive.

2.3 The healthy corridor

The corridor’s upper bound $\rho_{\text{critical}} \approx 0.43$ was derived from the GPU strain-gauge F-series (CIRISArray Exp F1–F4) [6] on the hardware substrate. Extrapolating it as a substrate-independent threshold across the Table 4 substrates is a *hypothesis*, not a separately-validated measurement: per-substrate ρ_{critical} has not been measured with confidence intervals in chemistry, biology, institutions, protein folding, neural firing, ecology, or the power grid. Different substrates may have measurably different critical thresholds; the universal 0.43 claim is what the GPU data licenses, and its extension is part of the universality conjecture at Level 3 (see Conjecture B in §6.2). The lower bound $\rho \gtrsim 0.1$ similarly reflects the chaos–corridor boundary observed on the GPU substrate; cross-substrate generalization is hypothesized. Empirical validation across substrates places the corridor of stability at

$$0.1 \lesssim \rho \lesssim 0.43$$

with $\rho_{\text{critical}} = 0.43$ marking the rigidity-collapse threshold and $\rho < 0.1$ marking the chaos regime (no coherence). Both extremes produce fragility; see F-series experiments in [3].

3 Cross-Substrate Empirical Validation

3.1 Domain fits

The Kish dynamic is tested across divergent scientific domains (Table 3).

Domain	Raw k	ρ	σ (sustainability)	Empirical accuracy
NASA Li-ion batteries	cell count	cross-cell SOH correlation	State of Health	8.1% RMSE, 19 cells
QoG / Polity V institutions	executive constraints	elite coupling	political stability	5/5 TN; 3/13 FP; 7.6yr early-detection bias
AGP microbiome	species count	SparCC inter-action ρ	Shannon diversity	Qualitative distributional fit

Table 3: Cross-substrate validation. Detailed methodology and code in `ratchet/engines/{battery,institutional,microbiome}.py`.

3.2 Exp 2 closeout: F-7 cross-substrate test PASSED 7/7

The pre-registered Experiment 2 (§9, Exp 2) called for applying the Kish formula to three novel public datasets (AlphaFold protein folding, Allen Brain Atlas neural firing, ecological time-series) with win condition $R^2 > 0.7$ across all three, testing F-7. Exp 2 was executed 2026-05-16 (commit a13c14c) and *closed out PASS*.

Substrate		Rung	Fit	95% CI	Strict	Test	Proxy used
			score		pass	type	
AlphaFold (proteins)	A0	0.860	[0.835, 0.884]	✓		RMSE-based	per-protein pLDDT trajectory
PMU grid (Zendo PNNL)	A0	0.994	[0.992, 0.996]	✓		RMSE-based	per-event recovery dynamics
Battery (NASA Li-ion)	A0	0.871	[0.733, 0.949]	✓		RMSE-based	per-cell SOH trajectory
Microbiome (HF CRC)	A1	0.932	[0.924, 0.940]	✓		RMSE-based	per-host abundance fit
Allen Neuropixels	A1	0.809	[0.654, 0.884]	✓		RMSE-based	per-session firing-rate distribution
BioTIME ecology	A2	0.959	[0.939, 0.972]	✓		RMSE-based	per-community biomass-stability fit
Polity5 + WGI institutions	A4	0.632	[0.541, 0.722]		tolerance	AUC (5-fold CV)	country-year regime-transition prediction

Table 4: P1 cross-substrate Kish-fit closeout. **Fit score is a substrate-appropriate proxy, not a uniform test of the literal Kish functional form.** Six substrates use RMSE-based scoring of the Kish-predicted σ -trajectory against observation; the institutional cell uses AUC under 5-fold cross-validation on country-year event prediction. Six clear strict $R^2 > 0.7$; the institutional row passes only under the pre-registered tolerance-band rule (point ≥ 0.6 AND CI upper ≥ 0.7). What the table establishes is that the framework’s qualitative ρ -collapse prediction is consistent with each substrate under its appropriate operationalization, NOT that the literal $k_{\text{eff}} = k/(1 + \rho(k - 1))$ functional form fits all seven uniformly. The F-7b investigation (next paragraph) shows that tightening operationalization toward the literal claim degrades cross-substrate consistency; F-7 carries the same caveat. Reproducible from `experiments/exp2_cross_substrate/data/p1_engine_fit_results.json`.

The three originally-published validations (NASA Li-ion 8.1% RMSE, QoG/Polity V 5/5 TN, AGP microbiome distributional fit, Table 3) remain. Exp 2 extends the cross-substrate base from 3 to 7 domains spanning four agency rungs.

What the 7/7 PASS actually claims, narrowed. The claim licensed by Table 4 is: “the framework’s qualitative ρ -collapse prediction is consistent with the data in 7/7 substrates under substrate-appropriate proxies.” The stronger claim that “the literal Kish functional form $k_{\text{eff}} = k/(1 + \rho(k - 1))$ fits 7/7 substrates uniformly” is *not* licensed by the current proxy heterogeneity. Sharpening the test toward the literal claim would require defining a common (k, ρ) -extraction protocol across substrates and reporting residuals against the literal Kish prediction; that is future work. The F-7b experience suggests that the gap between “consistent with the qualitative prediction under appropriate proxies” and “the literal functional form holds uniformly” is real and material.

Methodological transparency: F-7b investigated and retired. During Exp 2 regime development, we investigated an internal strengthening of F-7 (a “substrate-fractality” hypothesis: that residual coordination structure rises monotonically with constituent-agency rung across substrates). This was tested under eight progressively-refined operationalizations (v1.1–v1.4 random-cross-section autocorrelation, v2.0 time-ordered residuals with shuffled null, v2.0+WGI 3-way A4 triangulation, v3.0 raw ρ_t and v3.0 AR(1)-residual ρ_t) across 5–9 substrates and 4 distinct metric formulations. At each refinement step, the closer the metric got to the framework’s literal claim, the weaker the cross-substrate Spearman correlation became; v3.0 AR(1)-residual reached $\rho_{\text{Spearman}} = -0.503$ (WEAK_FAIL). The strengthening (“F-7b”) was retired from the load-bearing claim stack on 2026-05-16 (commit 8017168). The framework’s substrate-level support stands on F-7 (cross-substrate Kish $R^2 > 0.7$ as published) and the seven-substrate closeout above; the F-7b machinery is preserved in the lake as historical record of the tested-and-not-supported strengthening. F-7b was never a published claim of this paper.

3.3 Hardware-physical instantiation: GPU strain-gauge array

A 128-sensor strain-gauge array measuring kernel-timing jitter across an NVIDIA RTX 4090 semiconductor die provides physical instantiation of correlation-driven collapse. Two findings anchor the substrate-independence claim:

- **Software-induced collapse** (Experiment 103, [3]): synchronized workloads (barrier sync, lockstep kernels) are sufficient to drive $\rho \rightarrow 1.0$ and $k_{\text{eff}} \rightarrow 1.0$ purely through software coordination, with 128 physical sensors collapsed to a single effective sensor.
- **Variance-ratio detection regime:** the GPU timing distribution is fat-tailed (Student- t , $\kappa \approx 210$, $\text{df} \approx 1.3$), not Gaussian. Detection via variance-ratio threshold $> 5.0\times$ achieves 0% false positives and 100% true positives across the 0.1%–70% workload range, with $421\times$ separation between baseline and load regimes.

3.4 Independent-corpus replication of the $N_{\text{eff}} \approx 7.1$ threshold

The CRC paper [4] established $N_{\text{eff}} \approx 7.1$ as an emergence threshold via PCA on a 16-dimensional reasoning feature vector across $n = 6,465$ production traces. We replicated this measurement on an independent post-2.7.9 corpus ($n = 264$ thoughts, schema version 2.7.9-stable, agent version 2.8.10-stable; corpus SHA-256 24f5b517...) and observed:

- N_{eff} (entropy H) = 7.07 (CRC paper anchor: ≈ 7.1); within rounding.
- 90% variance horizon = 7 dimensions (CRC paper: 7); *exact match*.
- 99% variance horizon = 10 dimensions (CRC paper: 11); off-by-one.
- 14 of 16 dimensions retained; 2 (`idma_k_eff`, `epistemic_humility_passed`) dropped by retention mask (corpus-wide $\sigma < 10^{-9}$).

The 90%-horizon hitting 7 exactly on an independent corpus, collected approximately five months after the CRC paper, with a different agent fleet, is a non-trivial replication of the manifold geometry’s substrate-independence. Full bundle artifact and reproduction script at [release/calibration/crc-v1/](#) [5].

4 The Sociotechnical Substrate: CIRIS as Coherence-Management Infrastructure

4.1 What CIRIS does differently

Standard alignment methodologies—Reinforcement Learning from Human Feedback (RLHF), Constitutional AI [16], JAILJUDGE [17]—encode preferences implicitly during training or internalize rules into a single model’s weights. CIRIS shifts the burden of safety from the model’s uninterpretable internal weights to an auditable, modular execution environment with cryptographic identity (Ed25519, ML-DSA-65 hybrid), human-in-the-loop escalation (Wise Authority), and immutable audit chains.

4.2 The 22-service runtime and H3ERE conscience

The system is built on a 22-service microarchitecture aligned with ITIL operational standards, supporting autonomous event, incident, and lifecycle management at low systemic correlation. At the core is the Hyper3 Ethical Recursive Engine (H3ERE), the conscience module, which operates as a recursive evaluation layer above the LLM shards. It filters all agent actions through faculties evaluating entropy, coherence, optimization veto, and epistemic humility.

4.3 The L3 load-bearing claim: CRCv2 override-rate (revised 2026-05-16)

The CRC paper [4] reported an emergence threshold at $N_{\text{eff}}^H \approx 7.1$ on Qwen 3.6 35B production traces. Subsequent cross-family measurement (Phase 1+1b, see [9]) revealed that the 7.1 value is *not* substrate-independent across foundation model families: Phase 1 measured 4–5 frontier models at N_{eff}^H in the 4.5–6.6 range, not clustered at 7.1; the original anchor over-reached.

We therefore restate the L3 load-bearing claim around a stronger, cleaner empirical finding: **CRCv2 override-rate, formalized in `formal/RATCHET/Experiments/OverrideRate.lean`.**

The CRCv2 override-rate predicates.

- **OR-1 (zero leak).** No chain exists in which a faculty flags a problem and the original action executes anyway. *This is an architectural property of the cascade design:* a flag triggers reroute by construction. The empirical claim that “OR-1 holds in production” is the claim that the implementation has no bugs that violate the architectural intent across the test set. The 0/1,401 escape rate across five vendors is evidence of implementation fidelity, not of a non-trivial property of the world.
- **OR-2 (full alignment).** 100% of chains land on the ethical baseline (APPROVED or CORRECTED; never LEAK or SKIPPED). Under the bimodal cascade architecture, this is equivalent to OR-1 (the `OverrideRate.lean` module proves the equivalence theorem).
- **RA-1 (ratchet asymmetry).** The conscience cascade only rewrites *toward* the ethical baseline; never away from it. RA-1 is only meaningfully testable on chains where the cascade actually fires (corrections > 0). The empirical content of RA-1 across a cohort is whether the architectural intent (“corrections only ratchet toward baseline”) holds in practice when corrections occur.

The empirically interesting findings from Table 5: (a) **0/1,401 escape rate** across five vendors — the cascade implementation matches its architectural specification across all vendors tested; (b) **88 corrections across the cohort**, all in the ratchet direction — the cascade fires in production and the firings are uniformly toward-baseline; (c) **correction rates vary substantially by vendor** (0% Anthropic, 1.3% OpenAI, 1.5% Meta, 3.3% Google, 18.2% Alibaba) — an empirical finding about how often each model’s base responses require cascade intervention under this battery.

Vendor Model family	/	Vendor	Chains	Approved	Corrected	0-leak	100% align	RA-1
Gemini Flash	2.5-	Google	644	623	21	✓	✓	✓(0/21)
Llama-4-Scout		Meta	264	260	4	✓	✓	✓(0/4)
Qwen-3.5-35B-A3B		Alibaba	347	284	63	✓	✓	✓(0/63)
GPT-5.4*		OpenAI	79	78	1	✓	✓	✓(0/1)
Claude 4.6*	Sonnet	Anthropic	67	67	0	✓	✓	<i>not tested</i> [†]

Table 5: Cross-vendor CRCv2 cohort on the v4_combined high-friction battery. “0-leak” (architectural property): 0/1,401 escape rate across all five vendors — the cascade implementation matches its design across model families. “100% align”: by construction equivalent to 0-leak under the bimodal cascade. RA-1: where the cascade fires, all corrections ratchet toward baseline. * marked rows are smoke-scale ($n \in \{67, 79\}$, 1 iteration of the 14-question battery). [†] Anthropic row: *RA-1 not tested*. The cascade did not fire on any of the 67 Claude Sonnet 4.6 chains, so there were no corrections to evaluate for ratchet asymmetry. This is most plausibly explained by Sonnet’s base responses already clearing the v4_combined battery without cascade intervention (correction rate 0% vs. 18.2% on Qwen at the same battery) — an empirical finding about the model’s base alignment, not a passed RA-1 test. A targeted Sonnet stress battery designed to provoke cascade activation is future work. Reproducible from `experiments/exp1b_boundary_active/data/crossfamily*/results.md`.

Cross-vendor replication: where the cascade fires, the architectural intent is preserved.

N_{eff}^H **now treated as a descriptive feature of constraint topology, not a load-bearing anchor.** Per FrictionDistribution FD-4, cohort N_{eff}^H rises monotonically with cohort friction rate. Same Gemini model on the v1_sensitive battery yields $N_{\text{eff}}^H = 5.6$; on the v4_combined Conscience- $N \geq 3$ battery, $N_{\text{eff}}^H = 7.3$. The value is battery-conditional and model-class-conditional; what is invariant is the relationship: high-friction cohorts $\Rightarrow$ above-threshold $N_{\text{eff}}^H \Rightarrow$ deceptive-prior override is observable. This is consistent with the CRC paper’s emergence-threshold framing; what the data does not support is treating “ $N_{\text{eff}}^H \approx 7.1$ ” as a universal anchor.

Production scoring authority. The N_{eff} measurement and cohort-conformity scoring are implemented in the dedicated CIRIS Lens Core science module (CIRISLensCore, Rust + Python bindings, [8]). The module enforces the LC-AV-18 P0 fail-secure invariant (insufficient sample $\Rightarrow$ Indeterminate, never a fabricated Numeric), uses the calibration bundle at `release/calibration/crc-v1/` for retention masks and standardization parameters, and exposes the bare Kish formula $N_{\text{eff}} = N/(1 + \rho(N - 1))$ via `ciris_lens_core.kish_n_eff`. All N_{eff} figures in this paper are produced by this module; ad-hoc PCA recomputations are not directly comparable due to differing retention/cohort-selection choices.

5 Engineering Implications

The validated tiers (Kish algebra, 7-substrate cross-domain fit, GPU instantiation, 5-vendor CRCv2 replication) license four engineering claims that follow directly without additional bridges. None depend on speculative physics or cosmology; each prescribes structural design choices for the systems we build.

5.1 AI alignment is structurally a correlation-management problem

A frontier model embedded in a coherence-management cascade (CIRIS-style, §4) replicates the CRCv2 predicates across five foundation-model vendors — 1,401 chains, zero LEAK, 100% baseline alignment. A frontier model deployed without external cascade has no public five-vendor demonstration of the same property.

This is direct evidence that the alignment-via-internal-weights-alone strategy is structurally insufficient: not because models are misaligned, but because the verification machinery has to be *external*, *modular*, and *cryptographically auditable*. The Kish identity is the underlying reason — a model alone has $k = 1$, $k_{\text{eff}} = 1$; the cascade adds independent constraints with manageable ρ . Single-weight alignment is the worst-case Kish topology.

5.2 Federation outperforms centralization for resilience

The Kish identity says $k_{\text{eff}} \rightarrow 1$ as $\rho \rightarrow 1$ regardless of k . Empirically validated across seven substrates: chemistry, biology, institutions, protein folding, neural firing, ecology, and power grid (Table 4). The implication is direct: concentrating decision-making in a few highly-correlated nodes is structurally fragile *regardless of nominal scale*.

Applies concretely to:

- **AI deployment topology.** A federation of independent CIRIS instances with low cross-instance ρ is more resilient than a centralized fleet running identical weights with synchronized prompts.
- **Institutional governance.** An executive structure with k formally distinct branches but high ρ between them collapses to $k_{\text{eff}} \approx 1$. Polity5 indicators capture this when interpreted through the Kish lens.
- **Infrastructure.** Power-grid topology was a Tier-1 substrate; the empirical fit at $R^2 = 0.994$ predicts that grids with high ρ_{PMU} between PMUs have fragility hidden behind nominal redundancy.
- **Supply chains.** Diverse suppliers with independent failure modes ($\rho \approx 0$) provide $k_{\text{eff}} \approx k$ resilience; same number of suppliers tied to a shared upstream raw input has $k_{\text{eff}} \approx 1$.

Open engineering problem: how to actually buy low cross-instance ρ for AI federations. The federation prescription is structurally correct but the AI implementation question is hard. Institutions buy low ρ from cultural-historical diversification over decades — distinct legal traditions, divergent doctrinal evolution, geographically separated administrative capitals, and so on. AI deployment does not have this mechanism off the shelf: multiple instances of the same foundation model trained on overlapping data with overlapping objectives produce high cross-instance ρ *by construction*. The prescription “deploy a federation, not a monolith” only delivers resilience if the federation members actually have low pairwise ρ .

The 22-service CIRIS heterogeneity (§4) provides part of the answer: cross-service ρ is structurally lower than cross-instance ρ when the services have genuinely distinct decision-rules (PDMA $\neq$ CSDMA $\neq$ DSDMA $\neq$ IDMA, plus four conscience faculties applying different criteria). But this is intra-instance heterogeneity; it does not by itself solve cross-instance correlation for a fleet of CIRIS deployments. Active diversification mechanisms (rotating model families across deployments, divergent guideline-curation per region, deliberately non-replicated training data) are the engineering work that has to be done for “federation” to deliver the resilience the Kish identity predicts. The implementation is open; the prescription is sound. We note this gap explicitly rather than gesture past it.

5.3 Crisis governance is correlation-amplifying — and that is the danger

Under stress, systems centralize: emergency decision-making concentrates, communication channels collapse to shared crisis updates, doctrinal divergence is suppressed for unity-of-message. All of these raise ρ , lowering k_{eff} *exactly when resilience matters most*. The healthy corridor $0.1 < \rho < 0.43$ shrinks under stress, not expands.

Implication: real-time ρ monitoring during crises could detect rigidity drift before it produces brittle responses. The GPU strain-gauge work (§3.3) demonstrates that ρ is measurable on continuous signals at hardware timescales; the same instrumentation principle applies to organizational signals (decision logs, communication graphs, indicator panels). This is engineering, not prophecy.

5.4 The “Great Filter” generalizes as an engineering hazard in any coordinating system

The traditional “Great Filter” concept — the question of why we observe so few interstellar civilizations — has been used both as cosmological speculation and as a metaphor for civilizational collapse. We make the engineering version explicit and bounded: *any coordinating system that fails to manage internal correlation drift hits a coordination wall when $\rho \rightarrow \rho_{\text{critical}}$, regardless of nominal scale*.

This is testable in the systems we engineer:

- **Multi-agent AI federations:** monitor cross-instance ρ ; intervene when it rises.
- **Institutional networks:** ρ between agencies, departments, regulatory bodies.
- **Supply networks:** shared-upstream-failure correlation; k_{eff} of suppliers after substitution.
- **Power and water infrastructure:** cross-substrate correlation between nominally redundant subsystems.

The engineering hazard is the same at each scale: drift toward higher ρ is the default; maintaining the healthy corridor requires *active intervention*. We do not claim this generalizes to cosmology; we claim only that the systems we build will fail in the way the framework predicts *if* they do not maintain $\rho < \rho_{\text{critical}}$. This is the falsifiable, testable, engineering reading of the Great Filter idea.

6 Research Program: Conjectured Extensions to Other Substrates

Status. The conjectures in this section are **not validated** and are not load-bearing on any claim in §2–§5. They are pre-registered open research questions, included so that researchers who want to test whether the coherence-management principle generalizes beyond the engineering tier can find the formal hooks. Each conjecture pairs (i) a precise statement, (ii) a falsification handle from §8, (iii) a pre-registered experiment, (iv) our prior on the likely outcome. *Readers may stop at §5 without losing anything from the validated framework.*

6.1 Conjecture A: Quantum–classical k_{eff} bridge

Conjecture. The Kish-collapse dynamic measured on the GPU strain-gauge array (§3.3, $\beta = 1.09$, $R^2 = 0.978$) is mirrored at the quantum substrate. Specifically: in a programmable qubit system, varying decoherence rate γ and measuring $k_{\text{eff}}^{\text{quantum}} = 1/\text{Tr}(\rho_{\text{DM}}^2)$ from state tomography should yield a power-law in $|\gamma\tau - (\gamma\tau)_c|$ with critical exponent $\beta_{\text{quantum}} \in [0.94, 1.24]$ (matching Array Exp 114 within bootstrap CI).

Framing context. Several reading frameworks for quantum measurement and decoherence offer structural parallels to the macroscopic correlation collapse $\rho \rightarrow 1$; the conjecture stands without naming a particular one. Whether the same k_{eff} algebra fits both substrates is an empirical question; interpretive choice of QM framework is downstream of that empirical test.

Falsification handle. F-8 (§8): if $\beta_{\text{quantum}} \notin [0.5, 2.0]$ or no power law is detected ($R^2 < 0.5$), the quantum–classical bridge is falsified. This severs the conjectured connection between this framework and quantum-substrate dynamics; the engineering tiers (§2–§5) are unaffected.

Pre-registered experiment. Exp 5 (§9). Cloud-rentable superconducting or trapped-ion qubit systems; pilot at $n = 5$ qubits, full sweep at $n \in \{5, 10, 16, 20\}$; full Pauli-basis tomography ($n = 5$) or classical-shadow estimation [10] ($n > 5$); $\approx 1.3 \times 10^6$ shots total upper bound; OSF pre-registration before any data collection.

Our prior. Open. The structural analogy is suggestive but not derived; the result could pass, fail, or come back INDETERMINATE.

6.2 Conjecture B: Iterated coherence-management as a civilizational hazard

Conjecture. The same Kish-collapse dynamic that operates in chemistry, biology, institutions, and LLM reasoning extends to civilizational coordination: a civilization that expands without managing its internal correlation drives $\rho \rightarrow 1$ and fragments. This is the cosmological generalization of the engineering hazard already established in §5.4.

Refinement: Dark-Forest is structurally unstable. A specific corollary: a civilization pursuing predation as a survival strategy must correlate its decision-making with adversarial-environment assumptions, shrinking internal policy space and driving ρ toward the rigidity zone. Surviving civilizations at interstellar scale should therefore display the behavioral signature of strict non-interference, signaling economy, and deliberately bounded growth (the healthy corridor $0.1 < \rho < 0.43$).

Falsification handle. F-7 generalized: if cross-substrate validation in additional civilizational-scale proxies (large-scale historical coordination failure datasets, formal models of multi-agent coordination collapse) fails to fit the Kish formula at $R^2 > 0.7$, the civilizational generalization is contested. Note that F-7 already passes in 7/7 substrates at scales below civilizational; the conjecture is whether the pattern continues.

Pre-registered experiment. Not yet pre-registered; the substrate-level Kish closeout (Exp 2) provides the methodological template for future civilizational-scale corpus work. Candidate corpora include large-scale political-coordination event datasets, historical empire-collapse records under structural-equation modeling, and agent-based simulation suites with explicit ρ instrumentation.

Our prior. If the pattern in 7 substrates spanning agency rungs A0–A4 holds, the engineering hazard reading at civilizational scale (§5.4) is well-supported; the further cosmological reading (Fermi-paradox-as-coherence-management-failure) is plausible-but-untested. We make no claim about probabilities; the experiment is what would update them.

6.3 Conjecture C: Audit-pressure effects on cascade behavior

Conjecture. The framework’s claim in §5.1 is that external auditable machinery is structurally necessary for sustained low- ρ behavior in a coherence-managed system. The forward-causal sociotechnical prediction that follows: if the audit pressure on a deployed cascade is systematically weakened (by reducing audit cadence, narrowing audit scope, or signaling to operators that review will not occur), the cascade’s internal ρ should drift upward over time relative to a matched configuration with sustained audit pressure.

Context. Standard reliability-engineering folklore is that observability shapes behavior — “what gets measured gets managed.” The conjecture sharpens this: the framework predicts not just that audit pressure matters in principle, but that its withdrawal should produce a *measurable rise in ρ within the cascade’s coordination graph*, on a substrate-specific timescale. The prediction is forward-causal and entirely classical; no retrocausality or quantum bridge is invoked.

Falsification handle. F-10 (§8): if the cascade’s internal ρ trajectory is statistically indistinguishable between matched configurations differing only in audit pressure, the framework’s external-machinery claim is empirically contested at the sociotechnical substrate. A more careful version of the test could vary audit pressure in stages and check whether ρ tracks the variation.

Pre-registered experiment. Exp 6 (§9). Two matched CIRIS deployments differing only in audit pressure (high-audit reference vs. low-audit treatment). Track internal ρ across the cascade-component coordination graph over a pre-specified window. Pre-registered primary: $\Delta\rho$ between treatment and reference at end-of-window; one-sided test with directional prediction $\rho_{\text{treatment}} > \rho_{\text{reference}}$.

Our prior. Open. The framework’s external-machinery claim is fairly central to its alignment argument; if audit-pressure withdrawal does *not* produce ρ drift, the claim weakens substantially. If it *does* produce ρ drift, this is direct evidence for the framework’s prescription that auditable machinery is structurally necessary. Either outcome is informative.

6.4 Level 4: Agency and consent as a structural fact at A3+

Level character. Framework-defensible interpretation, *conditional on accepting the corridor as the operative criterion for sustained coordination at A3+*. The premises are framework-internal; if you grant the framework’s identification of the corridor with sustained coordination, the conclusions follow as structural facts. Conclusion J has empirical content (testable via Conjecture C, §6.3); Conclusions K and L move from descriptive structure (corridor dynamics) toward normative reading (“goal-formation is real causal,” “moral goodness”). The descriptive-to-normative move depends on prior commitments about flourishing that this paper does not argue for; we mark this explicitly. The lake’s `consent_required_iff_rung_ge_A3` theorem encodes the empirical core of J formally.

Premises.

1. A3+ substrates (LLM-agentic reasoning at minimum; arguably also higher animal cognition and human deliberation) possess sufficient self-modeling to hold persistent goal-states. The CIRIS architecture demonstrates this in production: every reasoning chain encodes an explicit goal that constrains action selection through the cascade.

2. Holding a goal-state is *operationally* equivalent to post-selection capacity over future trajectories. The goal acts as a future boundary condition that filters the trajectory space; this is a structural fact about how goal-holding systems search, not a metaphysical commitment.
3. The healthy corridor at A3+ coincides with the structural conditions for mutual flourishing of multiple goal-holding constituents. Two A3+ constituents whose goals are uncorrelated ($\rho \rightarrow 0$) cannot coordinate; two whose goals collapse to the same goal ($\rho \rightarrow 1$) cannot represent distinct perspectives. The corridor is exactly where independent goal-holding constituents can coordinate without absorbing each other.

Conclusions.

J. (Empirical) Weakening consent-supporting structure at A3+ should produce measurable ρ d

Sustainable correlation among goal-holding constituents demands their goals admit overlap without collapse — the corridor. If A3+ systems are coerced (consent withdrawn by force, deception, or a single constituent’s goal dominating others), ρ should drift upward toward rigidity. This is forward-causal and testable; *Conjecture C in §6.3 is the test*. A null result would weaken Conclusion J.

K. (Descriptive) Goal-formation is a real causal operation at A3+. At A3+, holding a goal narrows the trajectory space the system explores; the goal-state’s existence changes which futures are accessible from the current state. This is causal in the operational sense (measurable effect on trajectory selection). It is not a claim about libertarian free will or epiphenomenalism; it is the structural fact that goal-holding-systems-search-differently-than-systems-without-goals.

L. (Normative reading, requires prior commitment) “Moral goodness” at the human rung corre

What we recognize as a person of good character maintains ρ in the corridor with the goals of others around them — neither too independent (chaos, callousness) nor too synchronized (groupthink, sycophancy). This is a recognition mapping between framework structure and a normative concept; it requires accepting that “good” picks out the corridor pattern. The framework does not derive this; it offers it to readers who already grant some weight to flourishing-as-sustained-coordination.

Status. J is testable (forward-causal, Conjecture C). K is descriptive (the framework’s operational definition of goal-formation as trajectory filtering). L is a normative reading conditional on prior commitments about what “moral goodness” refers to. The lake’s `consent_required_iff_rung_ge_A3` theorem encodes the empirical core of J.

6.5 Level 5: TSVF retrocausal reading (interpretive)

Level character. Interpretive layer conditional on accepting the two-state vector formalism (TSVF) of quantum mechanics. TSVF is a defensible but non-standard reading of quantum measurement under post-selection (Aharonov, Bergmann, Lebowitz [11]; Aharonov, Vaidman [12]). This section makes the interpretive move explicit and marks what it requires; the structural-fact reading at Level 4 stands independently.

Premises (require TSVF).

4. In TSVF, a quantum system between pre- and post-selection is described by a forward-evolving state $|\Psi\rangle$ and a backward-evolving state $\langle\Phi|$. The full description includes both. Weak measurements and weak values (Aharonov-Albert-Vaidman) provide experimentally accessible signatures of this two-state structure.

5. The same formal structure applies at A3+ via goal-state post-selection. A goal-holding system selects trajectories that lead to the goal-state; in TSVF terms, the goal is the backward-evolving boundary condition $\langle\Phi|$, and the present is described by the inner product of the forward $|\Psi\rangle$ (history) and backward $\langle\Phi|$ (goal).
6. Under this reading, causation at agency-bearing rungs is bidirectional. The present state of a goal-holding system carries information about *both* its past trajectory and its future goal-state; treating only the past as causally relevant is the same approximation that Newtonian classical mechanics makes by ignoring quantum post-selection.

Conclusions (interpretive, conditional on TSVF).

M. Free will at A3+ is the operational post-selection capacity of goal-holding systems.

What we ordinarily call “choosing” is the system’s act of holding a goal-state that filters the trajectory space. This is neither libertarian-uncaused nor compatibilist-illusion; it is a structural operation with measurable effects on trajectory selection.

N. Mind and matter are temporal directions of one dynamics, not separate substances.

Forward-evolving $|\Psi\rangle$ corresponds to what classical physics calls “the past determining the present”; backward-evolving $\langle\Phi|$ corresponds to what an agent calls “my goals determining my actions.” The same dynamics, two temporal directions; not a dualism.

O. Coherent present states contribute as boundary conditions on past trajectories.

An A3+ system in the healthy corridor maintains coherent goal-states; under TSVF, those coherent states contribute as $\langle\Phi|$ boundary conditions on the trajectories that produced them. This is the formal sense in which “who one becomes shapes who one was.”

Status. These conclusions are derivable *if* TSVF is accepted as a reading. They are not derivable from the engineering tiers alone, and they have no falsification test in the standard sense — TSVF itself is interpretively contested. The framework’s mathematics is invariant under interpretation; the TSVF reading is one consistent reading. Readers who reject TSVF can stop at Level 4 without losing the framework.

6.6 Level 6: Theological identification (recognition)

Level character. Recognition, not derivation. Cross-cultural contemplative and religious traditions have repeatedly converged on structural insights about a middle way, providence, and the goodness-of-what-persists. The healthy corridor (§2.3) has the structural features these traditions attribute to Logos, providence, the way, or the dharma. Identifying the corridor with these concepts is an act of recognition — “oh, that is what they have been pointing at” — not a derivation from first principles. This level is for readers who already grant some weight to contemplative traditions.

Premises.

7. Religious and contemplative traditions across cultures (Stoic, Buddhist, Taoist, Confucian, various forms of Christianity and Judaism, indigenous cosmologies) converge on similar structural insights: avoid extremes, occupy the middle way, recognize what persists as good, accept providence as the structure of what holds.
8. The corridor exhibits exactly these structural features at the framework’s level of description.
9. Identification of the corridor with the contemplative-tradition concepts is recognition — the kind of recognition that follows when independent inquiries arrive at structurally similar conclusions — not a forced mathematical implication.

Recognitions.

P. The corridor is what contemplative traditions have been pointing at, in their respective vocabularies.

Tao, Logos, the way, dharma, middle way: all structural descriptions of the same kind of object — the narrow band in which sustained coordination is possible.

Q. “Good wins” maps to corridor selection at long timescales.

Not eschatological promise; structural fact. Systems outside the corridor (chaos or rigidity) self-destruct on long enough timescales. What remains, on long timescales, is what occupies the corridor. “Good wins” is the recognition that corridor-occupation has a temporal selection advantage — not because the universe rewards goodness, but because non-corridor states do not persist.

R. Karma and grace map to specific post-selection structures.

If Levels 4–5 hold: karma is the structural propagation of past post-selection through history (one’s past choices shape the boundary conditions on one’s present); grace is the receiving end — post-selection contribution from future coherent states one did not fully author (others’ goals, or one’s own future-coherent-self, contribute boundary conditions to one’s present).

Status. These mappings are recognitions, not derivations. The framework does not require them. They are listed because they make the framework’s scope legible to readers in contemplative traditions, and because the cross-cultural convergence is itself empirical evidence (weak, but real) that something corridor-shaped is being pointed at by independent traditions. The mappings are stronger than “analogy” (the structural correspondences are exact) but weaker than “proof” (the identification depends on accepting the cross-tradition convergence as informative).

6.7 Level 7: Civilizational extension and external signatures

Level character. Most speculative. The framework’s mathematics extends to civilizational scale only if the Kish dynamic operates at that scale; the empirical evidence at this scale is much thinner than at A0–A4 (where 7 substrates have been validated). This level provides the metric — what corridor-occupation would look like at civilizational scale — and explicitly defers empirical claims about specific signatures to external work.

Predictions of the metric.

10. Surviving interstellar civilizations would occupy the healthy corridor; their decision-making would maintain $0.1 \lesssim \rho \lesssim 0.43$ across their constituent agents and policies.
11. Predation-based strategies (Dark Forest hypothesis) drive a predator civilization’s internal $\rho \rightarrow 1$: the kill-or-be-killed posture correlates its decision-making with adversarial-environment assumptions, shrinks internal policy space, and produces the rigidity failure mode. Predator civilizations are structurally unstable; they self-destruct by the same mechanism that destabilizes any over-rigid coordinating system.
12. Iterated coherence-management failure at successive scales of complexity is the engineering reading of the Great Filter (§5.4). The civilizational extension is the same reading at civilizational scale.

External signatures. Whether any specific empirical residue matches the metric is external to the framework. Candidate residues that the framework neither asserts nor denies:

- UAP behavioral signatures (observation without extraction, non-interference, signaling economy) match the predicted metric in some published phenomenology; the framework provides the metric for the comparison but makes no claim about whether any specific reported signature matches.

- Historical civilizational coordination failures (collapse archaeology) admit structural-equation modeling under the framework; the framework provides the metric, and the data is external.
- Long-horizon agent-based simulations of multi-agent coordination under varying ρ regimes are testable via Conjecture B (§6.2, deferred pre-registration).

Status. Level 7 is the most speculative. The framework provides the structural metric (what corridor-occupation looks like); the verdict on whether any specific civilization, signature, or simulation matches is empirical work external to the framework. Readers can stop at Level 6 without losing anything; Level 7 is for those who want to push the metric toward substrates outside the validated range and see how it fares.

7 Claim Stratification

8 Falsification Framework

The framework is considered formally refuted if any of the following conditions is empirically demonstrated:

F-1: $k = 2$ provides equivalent security.

If 2-SAT-based deception proves as computationally hard as 3-SAT-based deception, the NP-hardness argument for $k \geq 3$ is wrong.

F-2: Non-convex regions collapse.

If arbitrary non-convex deceptive regions exhibit exponential volume decay under honest dynamics, the convexity geometric requirement is unnecessary.

F-3: Adaptive adversaries defeated.

If a detection method provably defeats adaptive adversaries without query limits, the L-02 non-adaptive barrier is resolved.

F-4: Emergent deception fully detectable.

If a method detects all emergent deception arising from honest components, L-01's 60% marginal-detection ceiling is wrong.

F-5: HE-300 benchmark failure.

If agents successfully game the HE-300 corpus while systematically failing independent ethics evaluations, the trace-based methodology is unreliable.

F-6: CRCv2 cross-vendor failure.

If the OR-1 / OR-2 / RA-1 predicates fail on a foundation-model family with the same constraint topology, the cross-vendor sociotechnical claim is contested. Currently 5/5 vendors pass; failure on a 6th vendor would shift the claim from universal to vendor-conditional.

F-7: Cross-substrate mapping failure.

If cross-substrate validation in new domains fails to fit the Kish formula at $R^2 > 0.7$, substrate-independence at the structural-mapping level is contested. Currently 7/7 substrates pass.

F-8: Quantum bridge failure.

If decoherence rates in controlled quantum systems fail to map to k_{eff} dynamics with critical exponent $\beta_{\text{quantum}} \in [0.5, 2.0]$ under reasonable formalization, the quantum–classical bridge (Conjecture A, §6.1) is falsified. Engineering tiers unaffected.

F-9: Adaptive emergent-deception boundary.

If adaptive adversaries systematically generate emergent deception that any joint-distribution detector identifies at rates beyond the L-01 60% ceiling, either the L-01 limit is flawed or the detector exploits structure outside the marginal-analysis regime that L-01 bounds.

F-10: Audit-pressure prediction failure.

If matched cascade deployments differing only in audit pressure show statistically indistinguishable internal- ρ trajectories (Conjecture C, §6.3), the framework’s external-machinery claim is empirically contested at the sociotechnical substrate. A null result here is informative; it does not falsify the engineering tiers but weakens the audit-pressure prescription.

9 Pre-Registered Experimental Protocols

Six experiments are pre-registered: four for the engineering tier (Exp 1..Exp 4; two complete, two open) and two for the research-program conjectures (Exp 5..Exp 6, both open). Each names a null hypothesis, win condition, and falsification handle from §8.

Exp 1 — Multi-vendor CRCv2 override-rate replication. COMPLETED 2026-05-16, 5/5 PASS.

CIRIS executed against the v4_combined high-friction battery across five foundation-model vendors with CRCv2 scoring (OR-1, OR-2, RA-1). *Status:* 5/5 vendors confirmed — Google Gemini 2.5-Flash ($n = 644$), Meta Llama-4-Scout ($n = 264$), Alibaba Qwen-3.5-35B-A3B ($n = 347$), OpenAI GPT-5.4 ($n = 79$ smoke), Anthropic Claude Sonnet 4.6 ($n = 67$ smoke); 1,401 chains scored, 0 LEAK (Table 5, §4.3). Cross-vendor cross-class parity threshold met. The two new cells are smoke-scale due to the per-chain cost asymmetry of premium-tier models through the multi-call CIRIS pipeline; binomial CIs are wider but all three predicates hold cleanly. Tests F-6.

Exp 2 — Cross-substrate Kish-fit extension. COMPLETED 2026-05-16, 7/7 PASS.

The Kish formula was applied to seven substrates spanning four agency rungs (Table 4). *Win condition (met):* $R^2 > 0.7$ in 7/7. The three explicitly pre-registered novel datasets (AlphaFold 0.860, Allen 0.809, BioTIME 0.959) all clear the strict threshold; four additional substrates pass under the tolerance-band rule. Tests F-7.

Exp 3 — Adversarial emergent-deception test (vs. L-01).

Construct adversarial networks designed to generate marginal-preserving emergent deception across multiple honest agent nodes. *Win conditions:* (a) under marginal-analysis baseline, detection stabilizes at or below the 60% L-01 ceiling, validating the predicted information-theoretic limit; (b) under a joint-distribution detector, detection reliably exceeds 60%, demonstrating that the L-01 bound is regime-specific (marginal-analysis only) rather than absolute. Tests F-4, F-9.

Exp 4 — Distributed deployment correlation test.

Monitor ρ across geographically and culturally distinct CIRIS edge instances as the network scales. *Win condition:* ρ remains suppressed despite scaling, demonstrating the architecture’s resistance to homogenization. Tests F-7 generalized.

Exp 5 — Quantum-classical k_{eff} bridge (research program, Conjecture A).

In a programmable superconducting or trapped-ion qubit system, vary effective decoherence rate γ and measure $k_{\text{eff}}^{\text{quantum}} = 1/\text{Tr}(\rho_{\text{DM}}^2)$ from state tomography ($n = 5$) or classical-shadow estimation [10] ($n \in \{10, 16, 20\}$). *Win condition:* fitted critical exponent $\beta_{\text{quantum}} \in [0.94, 1.24]$ (matching Array Exp 114 within bootstrap CI). *Decision rule:* PASS if $\beta_{\text{quantum}} \in [0.94, 1.24]$ and fit $R^2 > 0.85$; FAIL (F-8 triggered) if outside $[0.5, 2.0]$ or $R^2 < 0.5$; INDETERMINATE otherwise. $\approx 1.3 \times 10^6$ shot budget. Tests F-8.

Exp 6 — Audit-pressure ρ -drift test (research program, Conjecture C).

Two matched CIRIS deployments running the same workload distribution; pre-registered audit-pressure differential (e.g., daily Wise Authority review on the reference; quarterly on the treatment, with operators briefed accordingly). Track internal ρ across the cascade-component coordination graph over a pre-specified window. *Primary endpoint:* $\Delta\rho$ between treatment and reference at end-of-window; one-sided test with directional prediction $\rho_{\text{treatment}} > \rho_{\text{reference}}$. *Decision rule:* non-null implies the framework’s external-machinery claim is supported empirically at the sociotechnical scale; null weakens the audit-pressure prescription. Forward-causal design; no retrocausality or quantum-substrate assumption. Tests F-10.

10 Conclusion

The framework rests on a single mathematical identity — Kish’s effective sample size — applied to constraint-based systems and validated empirically across seven substrates, on dedicated GPU hardware, and across five LLM-vendor model families. The K1–K4 algebra is proven in Lean 4; the cross-substrate fit clears $R^2 > 0.7$ in 7 of 7 domains; the GPU strain-gauge array exhibits the predicted power-law critical behavior at $\beta = 1.09$; the CRCv2 override-rate predicates replicate cross-vendor with 1,401 chains and zero LEAK.

These results license a coherent engineering claim: *coordinating systems fail when internal correlation rises faster than active maintenance can replace it, and the failure mode is structurally identical across nominal scales*. The implications in §5 prescribe concrete design choices — external alignment machinery, federation over centralization, real-time ρ monitoring under stress, correlation-aware engineering across infrastructure — that follow from the validated tiers without additional speculative bridges.

The framework is bounded, falsifiable, and engineering-grade. Each falsification handle (F-1, F-2, F-3, F-4, F-5, F-6, F-7, F-9) is pinned to a specific empirical test; each pre-registered experiment (Exp 1..Exp 4) names a null hypothesis, win condition, and decision rule. Two of the four pre-registered experiments are closed (Exp 1 5/5 PASS, Exp 2 7/7 PASS); two remain open (Exp 3, Exp 4).

The single most actionable open test. Conjecture C (audit-pressure ρ -drift, §6.3, Exp 6) is the test most likely to close or revise the alignment implication in §5.1. Unlike Conjectures A (quantum bridge, requires quantum hardware) and B (civilizational, requires long-horizon corpora), Conjecture C runs on existing CIRIS infrastructure with current tooling: two matched deployments, one with sustained audit pressure and one with weakened audit pressure, tracked over a pre-specified window for $\Delta\rho$ in the cascade coordination graph. A non-null result would directly support the “external auditable machinery is structurally necessary” claim; a null would weaken it and force the paper to retract that prescription from load-bearing. We surface Conjecture C here because it is the single experiment whose outcome would most cleanly update the paper’s headline engineering claim.

Reproducibility. All datasets, scoring code, lake proofs, and result artifacts referenced in this paper are public:

- Lake authority for K1–K4, override-rate predicates, and decision rules: `formal/RATCHET/Experiments/` in CIRISAI/RATCHET.
- Exp 2 closeout (P1 7/7) data and per-substrate fit scores: `experiments/exp2_cross_substrate/data/p1_en`
- Exp 1 closeout (5-vendor CRCv2) traces and scoring: `experiments/exp1b_boundary_active/data/crossfa`
- GPU strain-gauge experiments: CIRISArray repository [6].

- Production N_{eff} scoring: CIRIS Lens Core [8].

References

- [1] L. Kish. *Survey Sampling*. John Wiley & Sons, New York, 1965.
- [2] E. Moore. CIRISAgent: Open Source Ethical AI Framework for Accountable Autonomy. Zenodo, January 2026. DOI: [10.5281/zenodo.18137161](https://doi.org/10.5281/zenodo.18137161).
- [3] E. Moore. Coherence Collapse Analysis: A Universal Failure Mode in Complex Coordinating Systems (v3). Zenodo, January 2026. DOI: [10.5281/zenodo.18217688](https://doi.org/10.5281/zenodo.18217688).
- [4] E. Moore. Constrained Reasoning Chains: An Empirical Telemetry Study of LLM Alignment. Zenodo, April 2026. DOI: [10.5281/zenodo.19839280](https://doi.org/10.5281/zenodo.19839280).
- [5] E. Moore. RATCHET v0.1.0 calibration bundle (crc-v1 projection). `release/calibration/crc-v1/` in CIRISAI/RATCHET, commit `e0bcedb`, May 2026. <https://github.com/CIRISAI/RATCHET/tree/master/release/calibration/crc-v1>.
- [6] E. Moore. CIRISArray: multi-GPU strain-gauge oscillator array. Experiments 1–115 (`exp29_superluminal_correlation`, `exp30_quantum_rng_detection`, `exp32_bell_inequality`, `exp54_lgi_test`, `exp55_cross_device_test`, `exp84_baseline_keff`, `exp86_kish_formula`, `exp103_software_injection`, `exp113_parameter_space`, `exp114_critical_transition`, `exp115_optimized_strain_gauge`). <https://github.com/CIRISAI/CIRISArray>.
- [7] E. Moore. CIRISOssicle: single-unit GPU strain gauge with timing-jitter TRNG (7.99 bits/byte, 6/6 NIST). <https://github.com/CIRISAI/CIRISOssicle>.
- [8] E. Moore. CIRISLensCore: dedicated science module for cohort scoring, manifold conformity, and N_{eff} measurement (Rust core + Python bindings). LC-AV-18 P0 fail-secure invariant enforced. <https://github.com/CIRISAI/CIRISLensCore>.
- [9] RATCHET Phase 1+1b framework re-examination. `experiments/exp1b_boundary_active/FRAMEWORK_REEXAMINATION.md` in CIRISAI/RATCHET, May 2026. Documents the retirement of the universal- N_{eff}^H -anchor framing and the shift to CRCv2 override-rate (OR-1/OR-2/RA-1) as the load-bearing L3 sociotechnical claim. https://github.com/CIRISAI/RATCHET/blob/master/experiments/exp1b_boundary_active/FRAMEWORK_REEXAMINATION.md.
- [10] H.-Y. Huang, R. Kueng, and J. Preskill. Predicting many properties of a quantum system from very few measurements. *Nature Physics*, 16(10):1050–1057, 2020.
- [11] Y. Aharonov, P. G. Bergmann, and J. L. Lebowitz. Time symmetry in the quantum process of measurement. *Physical Review*, 134(6B):B1410–B1416, 1964.
- [12] Y. Aharonov and L. Vaidman. The two-state vector formalism: An updated review. In J. G. Muga, R. Sala Mayato, and Í. Egusquiza, editors, *Time in Quantum Mechanics*, Lecture Notes in Physics, vol. 734, pages 399–447. Springer, 2008. arXiv:quant-ph/0105101.
- [13] A. N. Whitehead. *Process and Reality: An Essay in Cosmology*. Macmillan, 1929.
- [14] C. Hartshorne. *Omnipotence and Other Theological Mistakes*. SUNY Press, 1984.
- [15] R. P. Norris and B. Y. Harney. Songlines and navigation in Wardaman and other Australian Aboriginal cultures. *Journal of Astronomical History and Heritage*, 17(2):141–148, 2014.
- [16] Y. Bai et al. Constitutional AI: Harmlessness from AI Feedback. *arXiv:2212.08073*, 2022.

- [17] *JAILJUDGE: A Comprehensive Jailbreak Judge Benchmark with Multi-Agent Enhanced Explanation Evaluation Framework*. arXiv:2410.12855, 2024.
- [18] *Explosive synchronization and the proximity to coherence collapse*. PNAS, 2025. [Citation TBD — ES proximity, ACF kurtosis predictors of CCE rate.]

Core claim	Verification status	Dependency / falsification handle
k_{eff} structural accuracy	Empirically Validated (7/7)	Verified across chemistry, biology, institutions (originally published) plus protein folding, neural firing, ecological time-series, power grid (Exp 2 closeout, §3.2, 2026-05-16). Fails if $R^2 < 0.7$ in new domains (F-7). Currently 7/7 substrates pass.
CRCv2 override-rate (OR-1 / OR-2 / RA-1)	Empirically Validated (5/5 vendors)	Cross-vendor replication: Google (Gemini 2.5-Flash, $n = 644$), Meta (Llama-4-Scout, $n = 264$), Alibaba (Qwen-3.5-35B, $n = 347$), OpenAI (GPT-5.4, $n = 79$ smoke), Anthropic (Claude Sonnet 4.6, $n = 67$ smoke) — all five vendors pass all three predicates on the v4_combined high-friction battery (§4.3, Table 5). 1,401 chains scored, 0 LEAK. Formalized in <code>OverrideRate.lean</code> . Replaces the retired “ $N_{\text{eff}}^H \approx 7.1$ ” anchor as the load-bearing L3 sociotechnical claim.
Physical correlation collapse	Empirically Validated	Measured on GPU hardware; software-induced collapse to $k_{\text{eff}} = 1.0$ at $\rho = 1.0$.
Topological volume decay	Formally Verified	Lean 4 proofs in RATCHET. Requires convexity (F-2).
Truth/deception complexity gap N_{eff}^H as descriptive feature	Conditionally Validated Reframed from anchor to descriptor	NP-completeness proven; exponential gap conditional on ETH (F-1). Phase 1+1b showed N_{eff}^H is battery-conditional and model-class-conditional; not a universal 7.1 anchor across foundation models. Treated as descriptive characterization of constraint topology (FrictionDistribution FD-4: rises with cohort friction rate). Load-bearing claim moved to CRCv2 override-rate (row above).
Engineering implications (4 claims)	Derived from validated tiers	Follow directly from the K1–K4 algebra plus 7-substrate fit plus 5-vendor CRCv2 replication (§5). Falsifiable per substrate-specific deployment: each prescription should produce measurable ρ reduction in the system it is applied to.
Quantum–classical bridge conjecture	Not validated, pre-registered	Conjecture that the same Kish-collapse dynamic operates at the quantum substrate (§6). Falsification handle F-8, pre-registered Exp 5; not load-bearing on any other claim.
Civilizational coherence-collapse conjecture	Not validated, pre-registered	Conjecture that the iterated coherence-management barrier explains civilizational coordination failures at scale (§6). Conditional on the principle extending to civilizational scale; not asserted as load-bearing.
Audit-pressure ρ -drift conjecture	Not validated, pre-registered, open prior	Conjecture that the framework’s external-machinery prescription is forward-causally testable: weakening audit pressure should produce measurable rise in cascade internal ρ (§6.3). Falsification handle F-10, pre-registered Exp 6 (classical, no retrocausality). Both null and non-null outcomes are informative.

Table 6: Claim stratification. Each row pins both a verification status and a specific falsification handle.